

National University of Singapore
School of Computing
CS2105: Introduction to Computer Networks
Semester 1, 2018/2019

Assignment 2
Reliable File Transfer Protocol

Date Due: **21st October 2018, 23:59**

Overview

In this assignment, you will be writing the sending and receiving transport-level code for implementing a reliable file transfer protocol over UDP. The underlying channel is unreliable and may corrupt, drop or even re-order the packets at random. This assignment should be fun since your implementation will differ very little from what would be required in a real-world situation.

A Word of Advice

This assignment is potentially complex and time consuming. We advise you write your program incrementally and modularly. For example, start by dealing with packet corruption, then packet loss and so on. You should test your program after every major change and it will be useful to use a version control system like Git or Mercurial.

Pair Work

This assignment should be solved individually. However, if you are struggling, you may pair up with another student and work on the assignment together. The maximum number of students working together on a submission is 2. Pair work will incur a **2 mark penalty** for each student.

Under no circumstances should you solve it as a pair and then submit it as individual work. This is considered plagiarism and will be dealt with according to the university's protocol.

Plagiarism Warning

You are free to discuss this assignment with your friends. But, ultimately, you should write your own code. We employ a zero-tolerance policy against plagiarism. If a suspicious case is found, student would be asked to explain his/her code to the evaluator in person. Confirmed breach may result in zero mark for this assignment and further disciplinary action from the school.

Submission

You will submit your program on IVLE. Zip your files into a single zip file and name it “a2-<student number>.zip” and upload to the Assignment 2 folder. For example, if your student number is A0123456X, then you should name your file “a2-a0123456x.zip”.

It is your responsibility to ensure that your file is successfully uploaded into IVLE before the dateline. You should refresh the IVLE workbin and confirm that your file there.

The dateline for submission is at 21st October 2018, 23:59. Late submissions will have an additional penalty of 2 marks per day.

Special Instructions for Pair Work

If you are submitting this assignment as a pair, you are to include special comments at the top of your submitted source code as follow:

```
// PAIR WORK  
// A0123456X  
// A0654321X
```

where each lines begins with the comment character for the language, line 1 would be **PAIR WORK** and lines 2 and 3 are the student numbers of both students.

There will be no allowance for students who forget to declare that they solved the assignment as a pair. We advise you update the header of your source file as soon as you decide to work as a pair.

One and only one designated student should upload the zip file to IVLE. The other team member should never submit. It is up to you to decide who the designated submitter is.

Please do not change the designated submitter after submission. The same student should re-upload any update to the submission. Under no circumstance should the other student upload any zip file.

General Submission

Java

Your zip file should include two java files named **FileSender.java** and **FileReceiver.java**, which contains the classes **FileSender** and **FileReceiver** respectively. You may include other necessary java files in the zip file.

We will unzip the contents of your zip file and compile your submission using **javac *.java** and run each program as described in the problem. Java 9 on **sunfire** will be used. The location of the Java executables is at **/usr/local/java/jdk/bin/**.

Python

Your zip file should contain the files **FileSender.py** and **FileReceiver.py**. Any extra files that you use should be properly referenced and zipped in the proper directory structure.

We will use Python 3.7 on **sunfire** to run your code. The location of the Python executables on **sunfire** is at **/usr/local/Python-3.7/bin/python3**. Note this is NOT the default **python** that is on the path.

Architecture

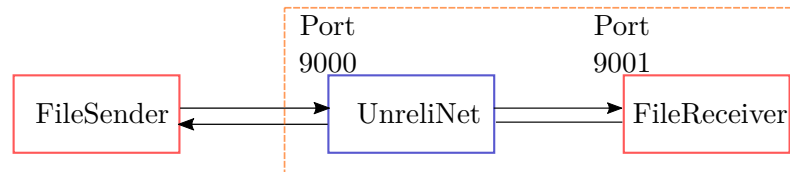


Figure 1: UnreliNET simulates an unreliable network.

There are three programs in this assignment: **FileSender**, **UnreliNET** and **FileReceiver**. The interactions of the programs is illustrated in Figure 1. The **FileSender** and **FileReceiver** programs implement a reliable file transfer protocol over UDP. The **UnreliNET** program acts as a proxy between **FileSender** and **FileReceiver**. Instead of the **FileSender** sending the packets directly to the **FileReceiver**, it sends all packets to **UnreliNET**. **UnreliNET** simulates an unreliable channel by randomly i) discarding packets, ii) corrupting packets and even iii) reordering the packets before forwarding them to the **FileReceiver**. Similarly, packets from the **FileReceiver** are also subject to the same unreliable conditions when forwarded to the **FileSender**.

The **UnreliNET** program is complete and given to you. Your task is to write the **FileSender** and **FileReceiver** programs so that a file can be successfully transferred in over an unreliable channel. The received file should be exactly identical to the one being sent. You will have to employ techniques taught in the lecture or from the textbook to ensure correct and reliable delivery.

Note: Protocols to handle out-of-order packets is not covered in the syllabus. Bonus marks will be given if your protocol is able to cater for this.

The FileSender

The **FileSender** program sends a given file to **FileReceiver** via **UnreliNET**. **UnreliNET** may then lose, corrupt or reorder the packets with a given probability before relaying them to the **FileReceiver** program.

The **FileSender** program will take in 4 input arguments: the host name, port number, source file, and the destination file name. For example:

```
$ java FileSender localhost 9000 ../large.mp4 big.mp4
```

sends the file `../large.mp4` to **UnreliNET** running on localhost on port 9000. **UnreliNET** will then relay the file to **FileReceiver** to be stored as `big.mp4`. The **FileSender** program will quit once the file has been successfully transferred.

You may assume that during testing, **FileSender** will be supplied with a valid path and filename, and that the pathname will not exceed 255 characters. No input validation is needed. You will also need to design a way to send the filename over to the **FileReceiver** for storing.

Note: Windows OS uses a different file separator '\', e.g., `..\large.mp4`

UnreliNET.class

We simulate an unreliable channel by relaying packets through the **UnreliNET** program. **UnreliNET** will only relay packets that contains at most 1,000 bytes of application data.

The compiled classfile of **UnreliNET** is provided. To run **UnreliNET**, the command is:

```
java UnreliNET <loss> <corrupt> <reorder> <unreliNET port> <recv port> [<random seed>]
```

where

- **loss** is the probability of losing a packet.
- **corrupt** is the probability of a packet being corrupted.
- **reorder** is the probability that a packet will be delayed and re-ordered.
- **unreliNET port** is the port that **UnreliNET** is to listen on.
- **recv port** is the port on localhost that the **FileReceiver** is listening on.
- **random seed** is an optional parameter which sets the random seed to obtain consistent behaviour for debugging. Setting the seed to 0 or omitting it will generate one using the system clock.

For example, the **UnreliNET** command for the setup shown in Figure 1 could be:

```
$ java UnreliNET 0.1 0.2 0.3 9000 9001
```

where on average 10% of the packets will be lost, 20% of the packets will be corrupted, and 30% of the packets will be delayed and potentially delivered out-of-order.

We recommend setting these values between the range [0.0, 0.3] for testing. Setting too large a value may result in very slow file transmission. If you have trouble getting your code to work, you might want to set them to 0 first for debugging.

You may also store the output of **UnreliNET** to a file for analysis as so:

```
$ java UnreliNET 0.1 0.2 0.3 9000 9001 > output.txt
```

FileReceiver.class

The **FileReceiver** program receives a file from **FileSender** (via **UnreliNET**) and saves it in the same directory it is running in, with the filename specified by the **FileSender** program.

It takes in one argument, the port number of which to listen on, e.g.,

```
$ java FileReceiver 9001
```

will start the process listening on port 9001.

The **FileReceiver** does not need to detect the end of transmission and terminate and can run indefinitely. This simplifies implementation otherwise you would have to implement a connection teardown like TCP. You may also assume that after receiving a file, the **FileReceiver** will not be reused for another file transfer.

<Ctrl>+C (SIGINT) will be sent to the **FileReceiver** once the **FileSender** terminates, so you should ensure that the file is correctly written to disk at this point in time. One way to ensure this is to flush the file buffer and close the file stream properly once all the data is successfully received.

Details and Grading

The score for this assignment is 10 marks. If you obtain above 10 marks, 50% of the extra marks can be used as bonus top-up for the assignment CA component.

A time limit will be imposed depending on the file size and error rate. More details after we perform some testing on **sunfire**.

Basic requirements

- 1 mark:** Your programs can be compiled on `sunfire` and executed without errors.
- 1 mark:** Your programs can successfully transfer a small file (a few KB) from sender to receiver over a perfectly reliable channel, i.e., no error is introduced.
- 1 mark:** Your programs can successfully transfer a large file (< 4 GB) from sender to receiver over a perfectly reliable channel, i.e., no error is introduced.
- 2 marks:** Your program can successfully transfer a file (both small and large) from sender to receiver in the presence of packet corruption.
- 2 marks:** Your program can successfully transfer a file (both small and large) from sender to receiver in the presence of packet loss.
- 2 marks:** Your program can successfully transfer a file (both small and large) from sender to receiver in the presence of both packet corruption and packet loss.
- 1 mark: [Speed test]** Your program can successfully transfer a large file (≈ 50 MB) from sender to receiver in the presence of both 10% packet corruption and 10% packet loss within 180 s.¹

Advance requirements

The following requirements can be quite challenging and can give you bonus marks.

- 1 mark:** Your program can successfully transfer a file from sender to receiver in the presence of packets being reordered. This is not covered in the syllabus so you will have to do your own research.
- 1 mark:** Your program can successfully transfer a file from sender to receiver in the presence of packet corruption, packet loss and reordering.
- 2–4 marks: [Speed Competition]** It is not difficult to ensure reliable delivery using the simple alternate bit protocols we discussed in lecture. However, they have very poor utilization. Thus, for this test, you will need to implement pipelining to increase the speed of the transfer. We will rank you according to the time your program takes to transfer a large file (≈ 50 MB) in the presence of 2% packet corruption and 2% packet loss.
 - Top 20 or within 110% of fastest time. **(4 marks)**
 - Top 50 or within 125% of fastest time. **(3 marks)**
 - All others which time is under 60 s.¹ **(2 marks)**

Note: Printing output to the terminal can cause additional delay. You may disable the output of `UnreliNET` by redirecting it to `/dev/null`, e.g., “`java UnreliNET 0.02 0.02 0 8000 8001 > /dev/null`”.

¹File size and time limit subject to change while we perform some testing on `sunfire`

Testing

Although **UnreliNET** is designed to accept traffic over a network, you should run all three programs on the same host to avoid interference from network traffic.

You should first run **FileReceiver** and **UnreliNET** before running **FileSender**. When testing on **sunfire**, you can open three SSH windows, one for each program. **UnreliNET** runs indefinitely and you may reuse it to run consecutive transfers. Press **<Ctrl>+C** to terminate it.

To check if both files are identical, you can use the **cmp** command just like in assignment 0 e.g., **cmp file1 file2**. If no output is produced, it means the files are identical.

Additional Tips and Pointers

Timer and Timeout Values

We do not recommend that you use a timeout value larger than 100 ms. Using a larger value might cause your program to be slow in sending data and exceed the allowable time limit. For simple alternating-bit protocols, you could potentially use a 1 ms timeout when testing on localhost as the RTT is below 1 ms.

Computing Checksum

UnreliNET will corrupt the payload of the UDP packet. Thus, you should use a checksum to verify the integrity of the received packets. You can use the same method as in Assignment 0 for computing a CRC32 checksum.

Self-defined Header/Trailer Fields

As part of your protocol, you will have to implement certain header/trailer fields into the packets, such as sequence number and checksum. The number of fields and the structure in a packet is agreed upon between the sender and the receiver. That is an application layer protocol that you design.

Note that each packet sent should contain at most 1,000 bytes of application data, inclusive of user-defined header/trailer fields. Otherwise, **UnreliNET** will drop the packet.

The **ByteBuffer** class from the **java.nio** package makes it easy to insert and retrieve primitive data into a byte array. For Python you can make use of the **structs** package. This is just a suggestion and you do not need to use this class if you have an alternative method.