

# CS2105 Introduction to Computer Networks

## Lecture 5

### UDP and TCP

10 September 2018

# Midterm Test

Monday 8 October 2018

- Week 8
- Time: 2:15pm – 3:15pm
- Venue: MPSH 2

## MRQs

- Multiple Response Question
- More fun than MCQ
- Partial score → higher random expectation

## Short Qns?

# Designing Reliable Protocols

Network layer service is unreliable.


- Packets may be lost or corrupted during transmission.

A reliable transport protocol should

- ensure that packets are received by receiver in good order.

Scenario: one sender, one receiver

- Sender sends data packets to receiver
- Receiver feeds back to sender

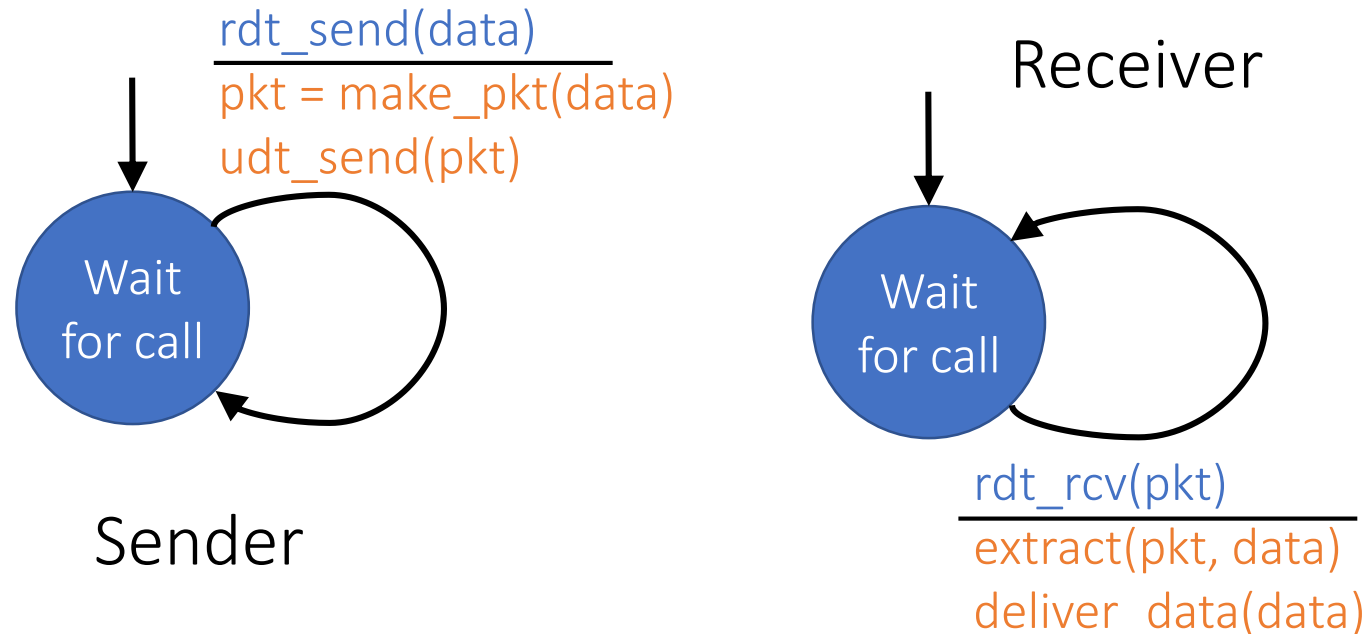


# rdt 1.0: Reliable Channel

Assume underlying channel is perfectly reliable.

No error checking is needed.

- Sender sends data into underlying (perfect) channel
- Receiver reads data from underlying (perfect) channel



# rdt 2.0: Channel with Bit Errors

## Assumption:

- underlying channel may flip bits in packets
- other than that, the channel is perfect

Receiver may use **checksum** to detect bit errors

Question: how to recover from bit errors?

- **Acknowledgements (ACKs)**: receiver explicitly tells sender that packet received is OK.
- **Negative acknowledgements (NAKs)**: receiver explicitly tells sender that packet has errors.
- Sender retransmits packet on receipt of NAK.

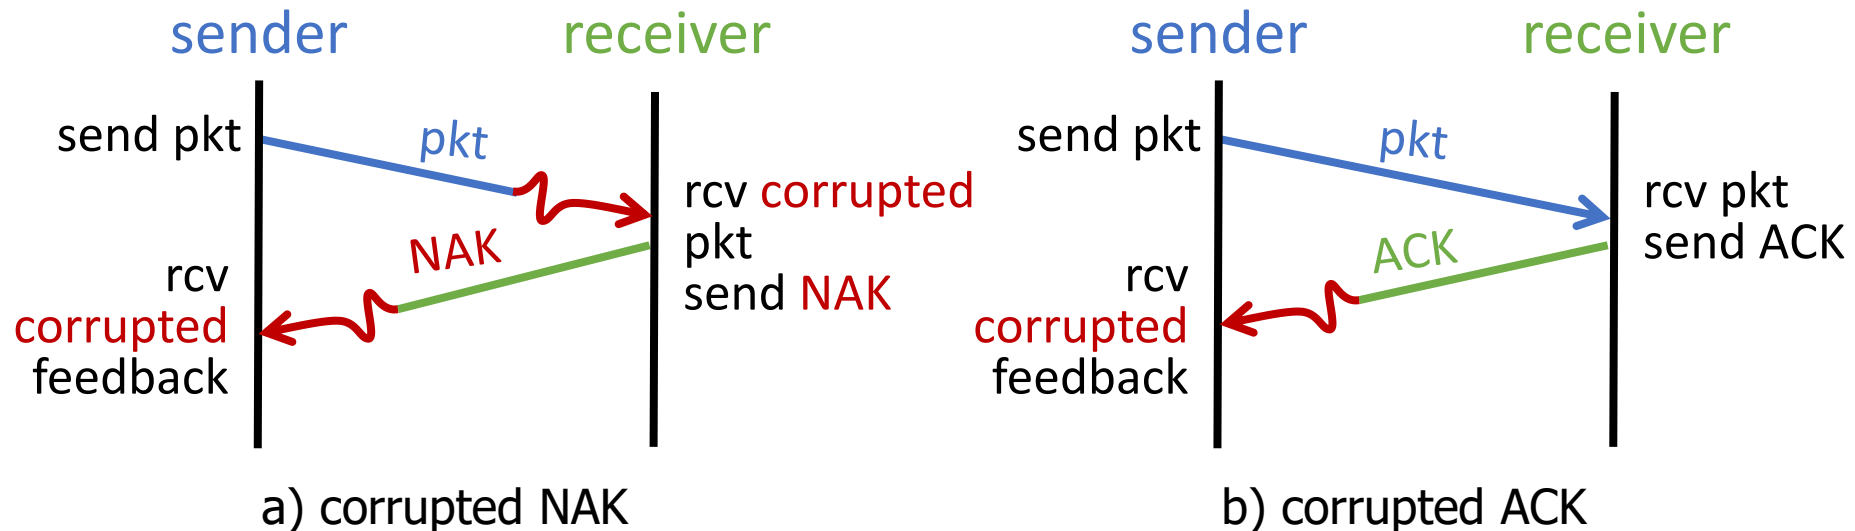
# rdt 2.0 has a Fatal Flaw!

# What happens if ACK/NAK is corrupted?

- Sender doesn't know what happened at receiver!

## So what should the sender do?

- Sender just retransmits when receives garbled ACK or NAK.
- Question: does this work?

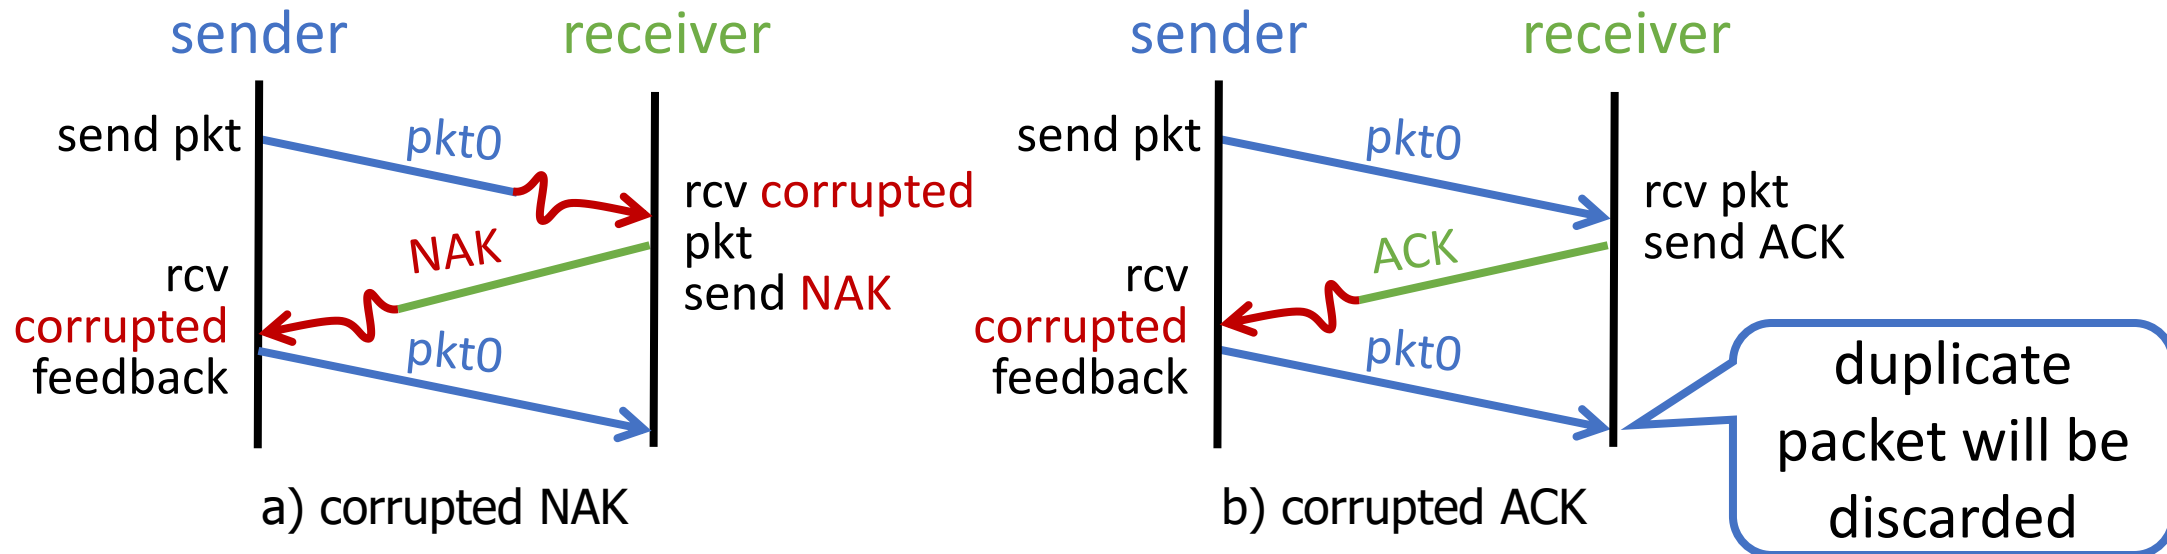


# rdt 2.1: Add a sequence number

To handle duplicates:

- Sender retransmits current packet if ACK/NAK is garbled.
- Sender adds sequence number to each packet.
- Receiver discards (doesn't deliver up) duplicate packet.

This gives rise to protocol rdt 2.1.



# rdt 2.2: a NAK-free Protocol

Same assumption and functionality as rdt 2.1, but use ACKs only.

Instead of sending NAK, receiver sends ACK for the last packet received correctly.

- Now receiver must **explicitly** include seq. # of the packet being ACKed.

Duplicate ACKs at sender results in same action as NAK:  
retransmit current pkt



# rdt 3.0: Channel with Errors and Loss

Assumption: underlying channel

- may flip bits in packets
- may lose packets
- may incur arbitrarily long packet delay
- but will not re-order packets

To handle packet loss:

- Sender waits “reasonable” amount of time for ACK.
- Sender retransmits if no ACK is received till timeout.

Sender relies on timeout/retransmission to deal with both packet corruption and packet loss.

# Learning Outcomes

After this class, you are expected to:

- know the simplicity of UDP and the service it provides.
- know how to calculate the checksum of a packet.
- know the operation of the components of TCP
  - sequence number
  - acknowledgement number
  - retransmission,
  - receiver window
  - connection setup and termination

# Chapter 3: Roadmap


3.1 Transport-layer services

3.2 Multiplexing and de-multiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP




You can't imagine how **simple**  
**UDP** is.

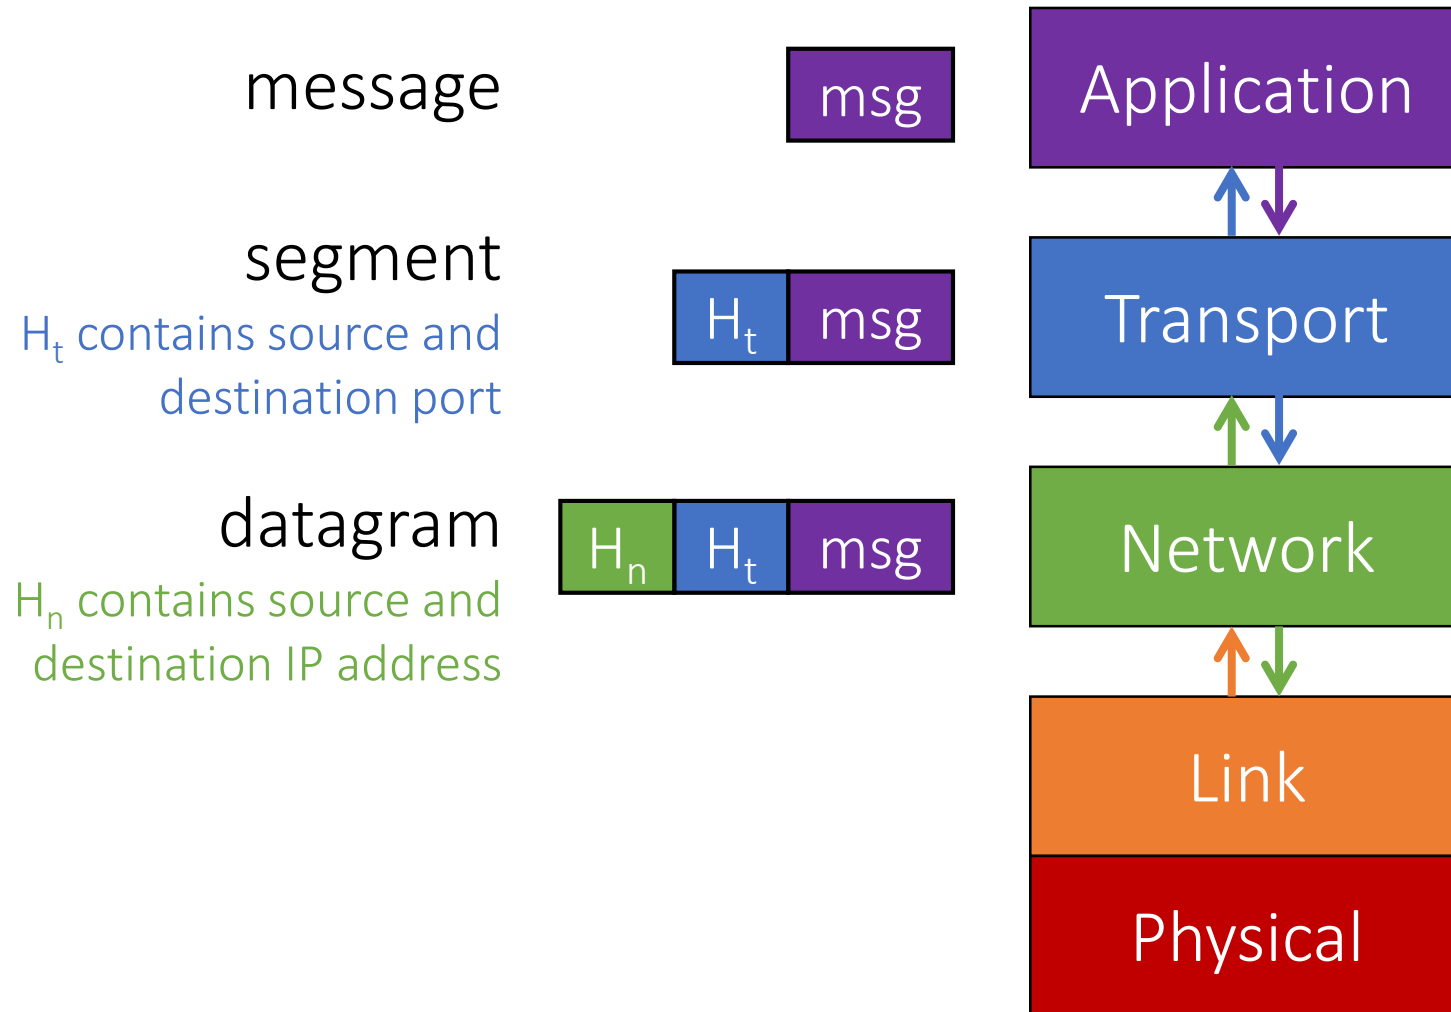


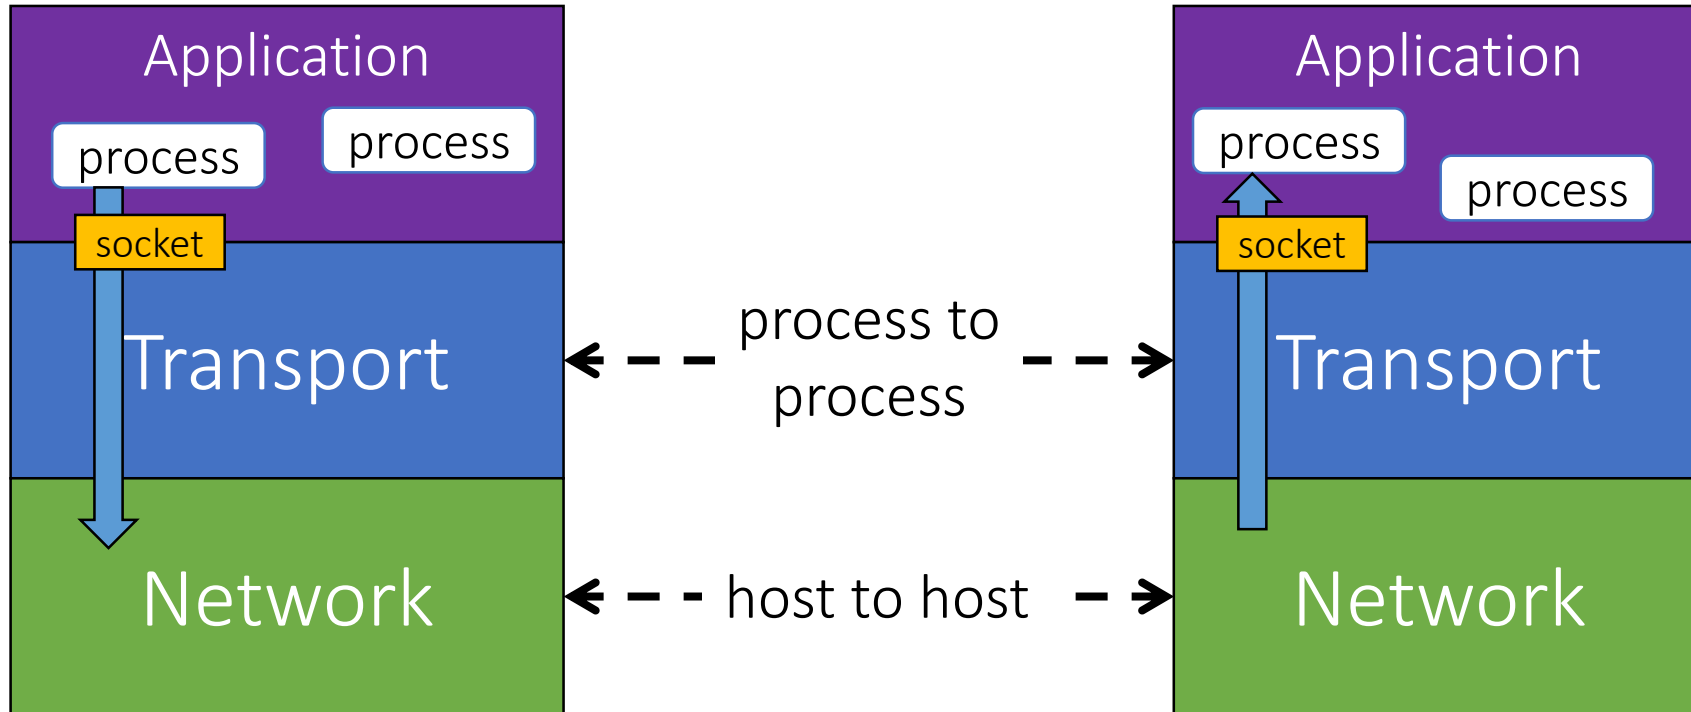
**IT'S SO SIMPLE**

But the complexity of TCP will  
make you cry



# Layering





# User Datagram Protocol

[RFC 768]



Internet protocols are described in documents  
known as Request for Comments (RFC)

<http://www.ietf.org/rfc/rfc768.txt> is only 3 pages

# UDP

adds very little service on top of IP:

- Connectionless multiplexing / de-multiplexing
- Checksum

transmission is unreliable

- Often used by streaming multimedia apps (loss tolerant & rate sensitive)

To achieve reliable transmission over UDP

- Application implements error detection and recovery mechanisms!

# Connectionless De-multiplexing

## UDP sender:

- Creates a socket with local port #
- When creating a datagram to send to UDP socket, sender must specify destination IP address and port #

## When UDP receiver receives a UDP segment:

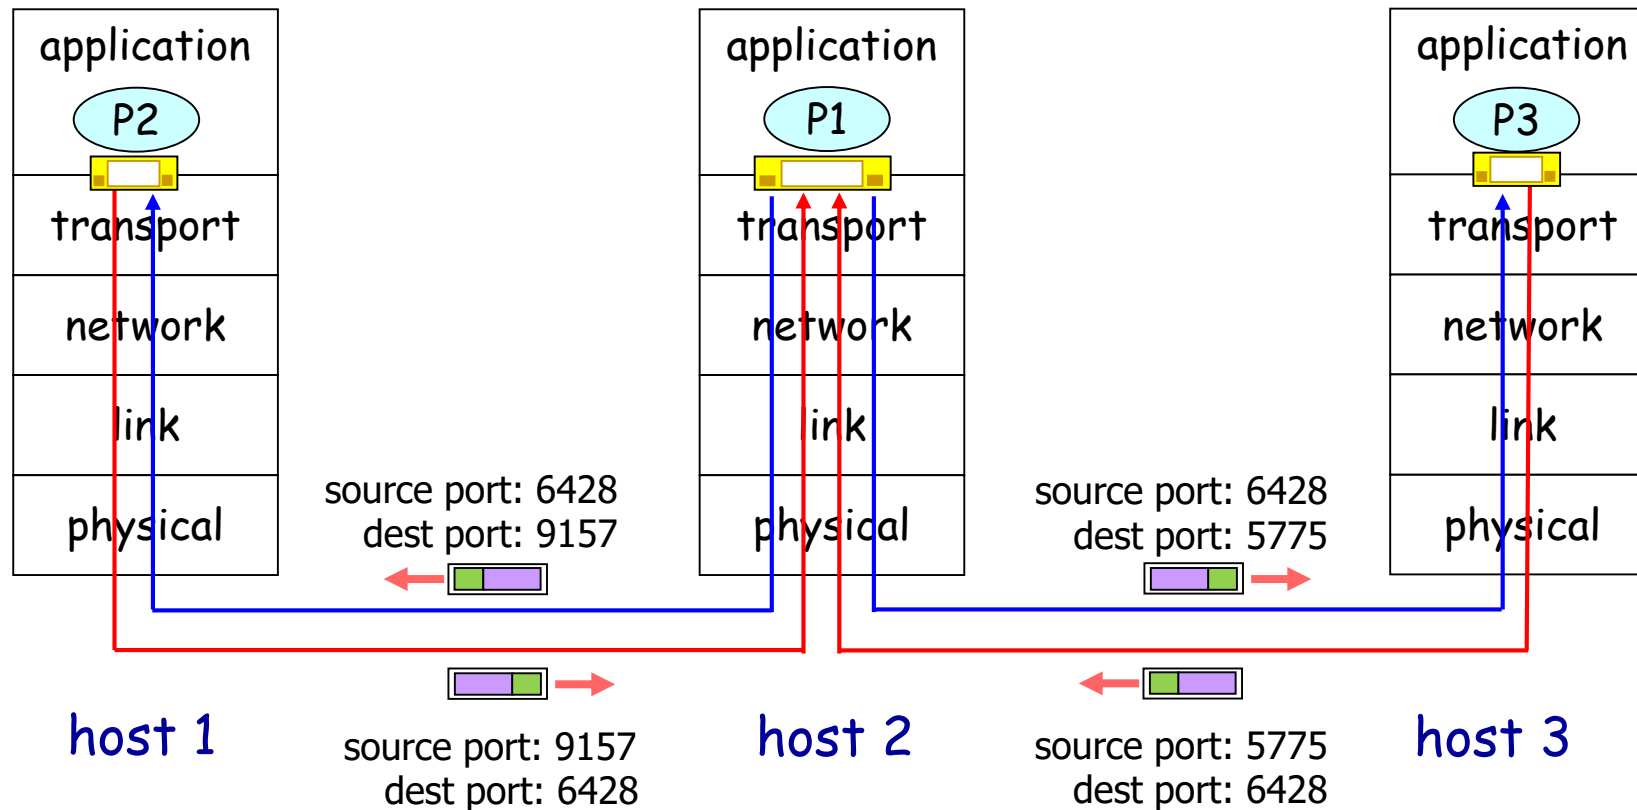
- Checks destination port # in segment
- Directs UDP segment to the socket with that port #
- IP datagrams (from different sources) with the same destination port # will be directed to the same UDP socket at destination

# Connectionless De-multiplexing

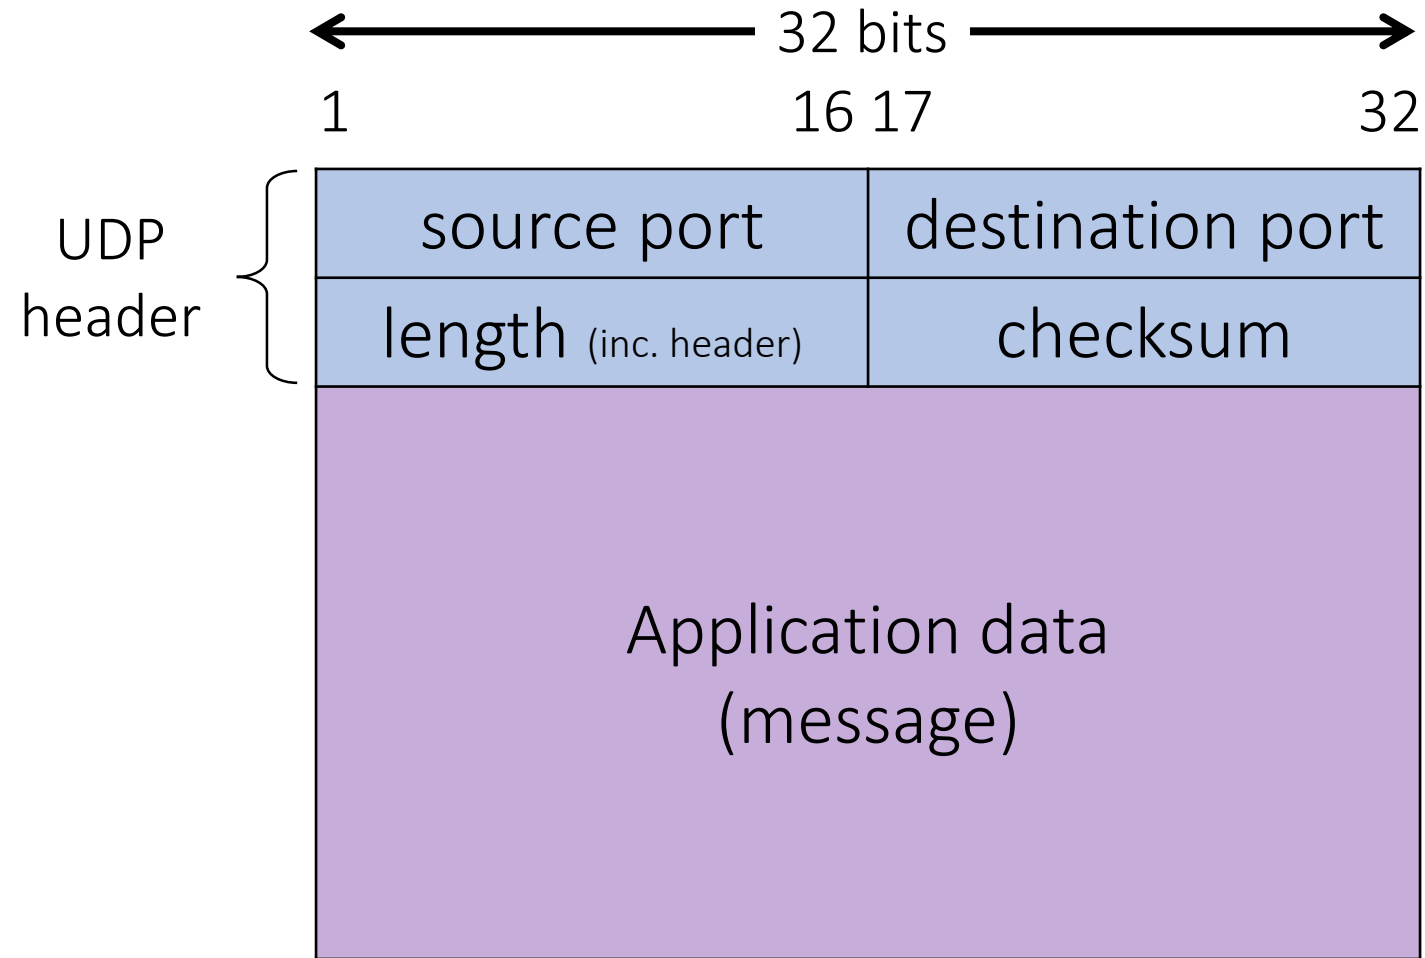
```
DatagramSocket mySocket2 =  
    new DatagramSocket(9157);
```

```
DatagramSocket mySocket1 =  
    new DatagramSocket(6428);
```

```
DatagramSocket mySocket3 =  
    new DatagramSocket(5775);
```



# UDP segment structure



# Why use UDP?

No connection set-up

- Reduce delay

No connection state at sender or receiver

- Need less resources

Small header size

- Less overhead

No congestion control

- Can blast as fast as desired

# Checksum

## Several different checksum algorithm

- Cyclic Redundancy Check (CRC)
- Message Digest v5 (MD5)
- Secure Hash Algorithm 1 (SHA-1)
- Secure Hash Algorithm 2 (SHA-2)
- UDP/TCP Checksum (RFC 768)

# UDP Checksum

Purpose: to detect errors (single bit flips) in the transmitted segment.

## Sender

Compute checksum value

- $f(P_s) = c_s$

Include checksum value into UDP checksum field

## Receiver

Compute checksum value

- $f(P_r) = c_r$

Compare if computed checksum equals checksum field value:

- NO - error detected
- YES - no error detected (but really no error?)




# Computing Checksum

1. Split segment into 16-bit integers (checksum initially 0)

```
1110 0110 0110 0110
1101 0101 0101 0101
```

# Computing Checksum

1. Split segment into 16-bit integers
2. Add next integer with wrap around carry



# Computing Checksum

1. Split segment into 16-bit integers
2. Add next integer with wrap around carry
3. Compute 1's complement

	1 1 1 0	0 1 1 0	0 1 1 0	0 1 1 0
	1 1 0 1	0 1 0 1	0 1 0 1	0 1 0 1
1	1 0 1 1	1 0 1 1	1 0 1 1	1 0 1 1
	<hr/>			
				1
sum:	1 0 1 1	1 0 1 1	1 0 1 1	1 1 0 0
checksum:	0 1 0 0	0 1 0 0	0 1 0 0	0 0 1 1

# Intuition

Using base10 integers

- Ensure no change to a sequence of numbers:  
27, 48, 13, 73, 52

Compute the sum

- $27 + 48 + 13 + 73 + 52 = 186$

But what if you only have 2 digits?

- Truncate? 86
- Wrap around?  $86 + 1 = 87$

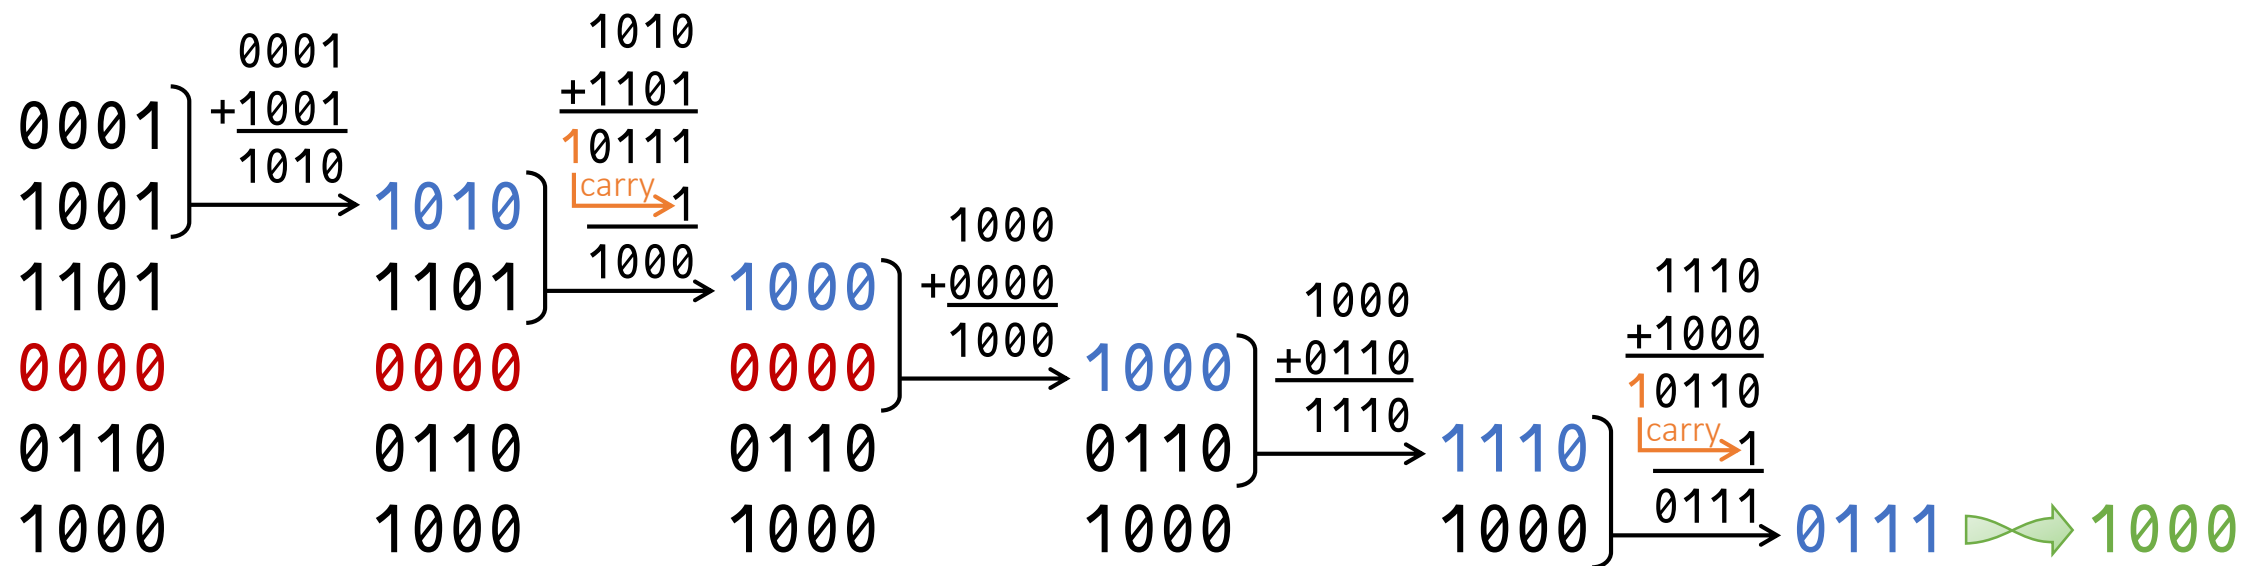
Take compliment?

- $87 \rightarrow 12$

# Example

Using 4-bit integers for simplicity

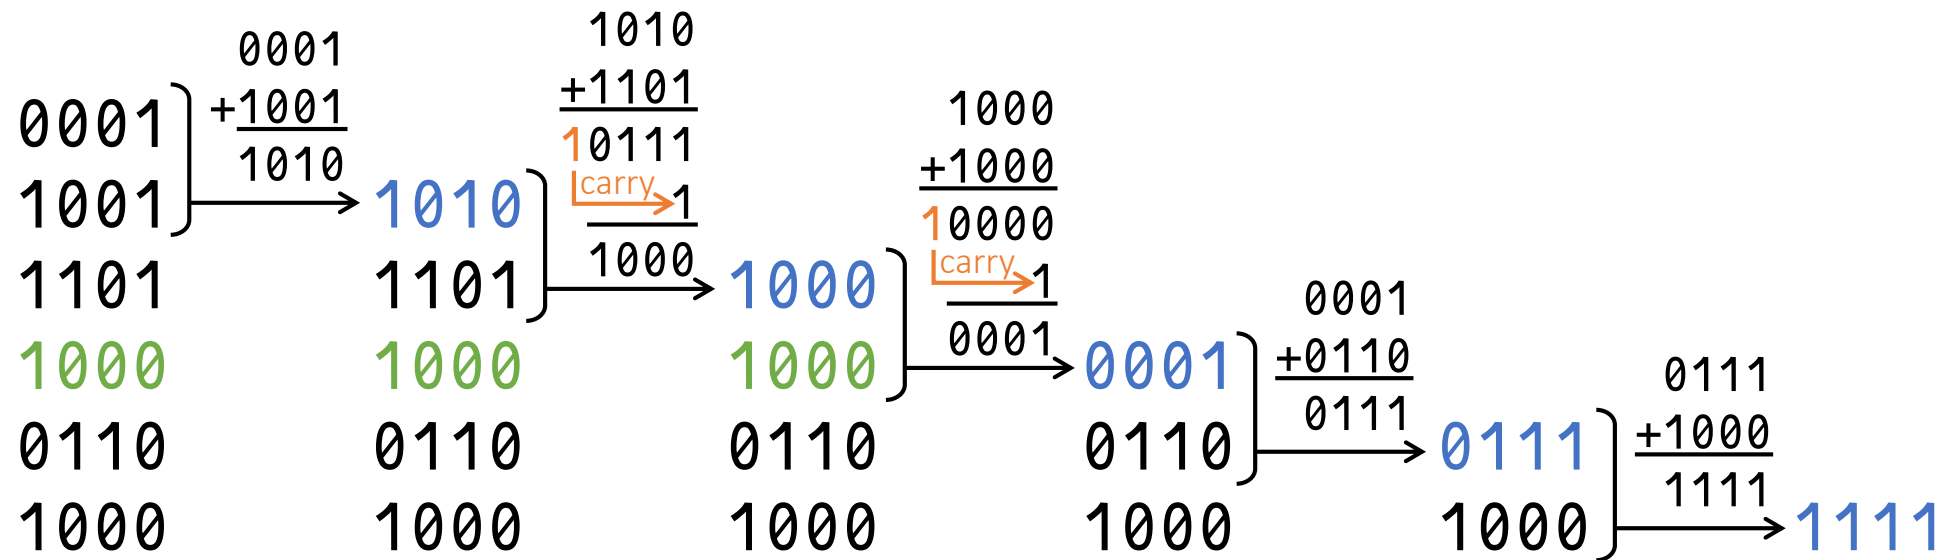
- Sender computes checksum for this message



# Example

Using 4-bit integers for simplicity

- Receiver verifies checksum for this message



# Chapter 3: Roadmap


3.1 Transport-layer services

3.2 Multiplexing and de-multiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP



# Transmission Control Protocol

[RFC 793, RFC 1122, RFC 1323, RFC 2018, RFC 2581, ...]

We only scratch the surface of TCP in CS2105.  
More will be covered in CS3103.



# TCP Overview

## Connection-oriented

- handshaking (exchange of control messages) before sending app data

## Reliable, in-order byte stream

- Application passes data to TCP and TCP forms packets in view of **MSS (maximum segment size)**
- (For UDP, app forms packets: **DatagramPacket**)

## Flow control and congestion control

- Not discussed!

# Re-cap

	Go-Back-N	Selective Repeat
ACK	cumulative	selective
out-of-order	ignore	buffer
retransmit	all unACK	one unACK
timer	oldest unACK	one per unACK


Which is TCP?

You should be able to answer after this lecture

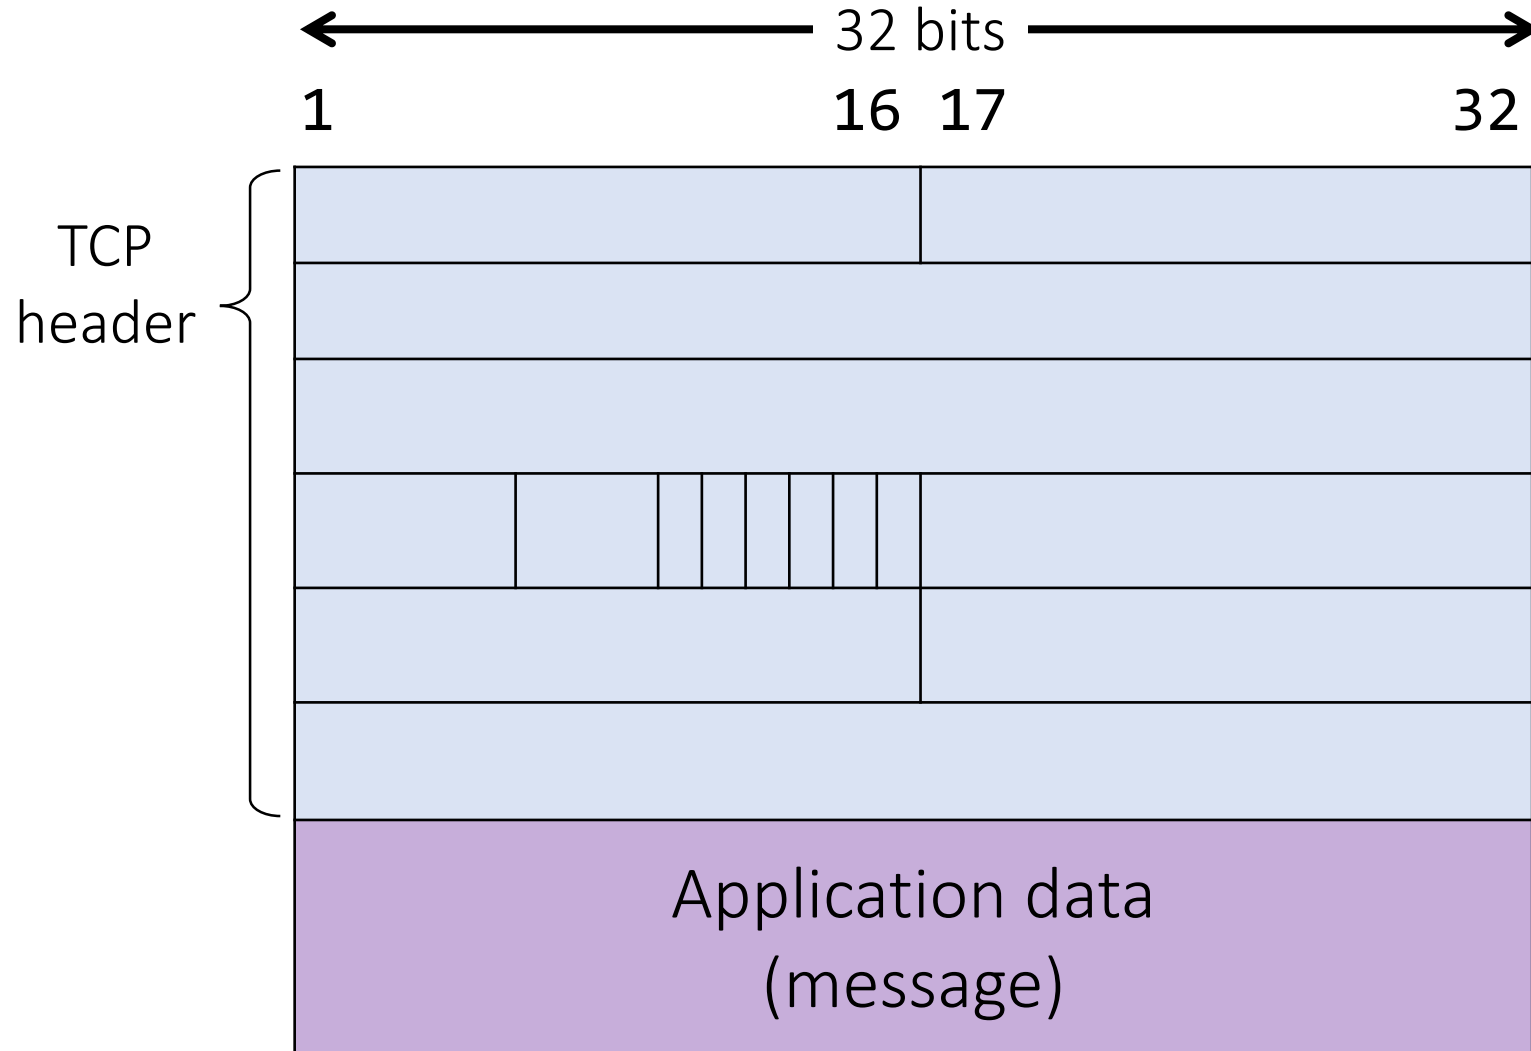
# Connection-oriented De-mux

A TCP connection (socket) is identified by 4-tuple:

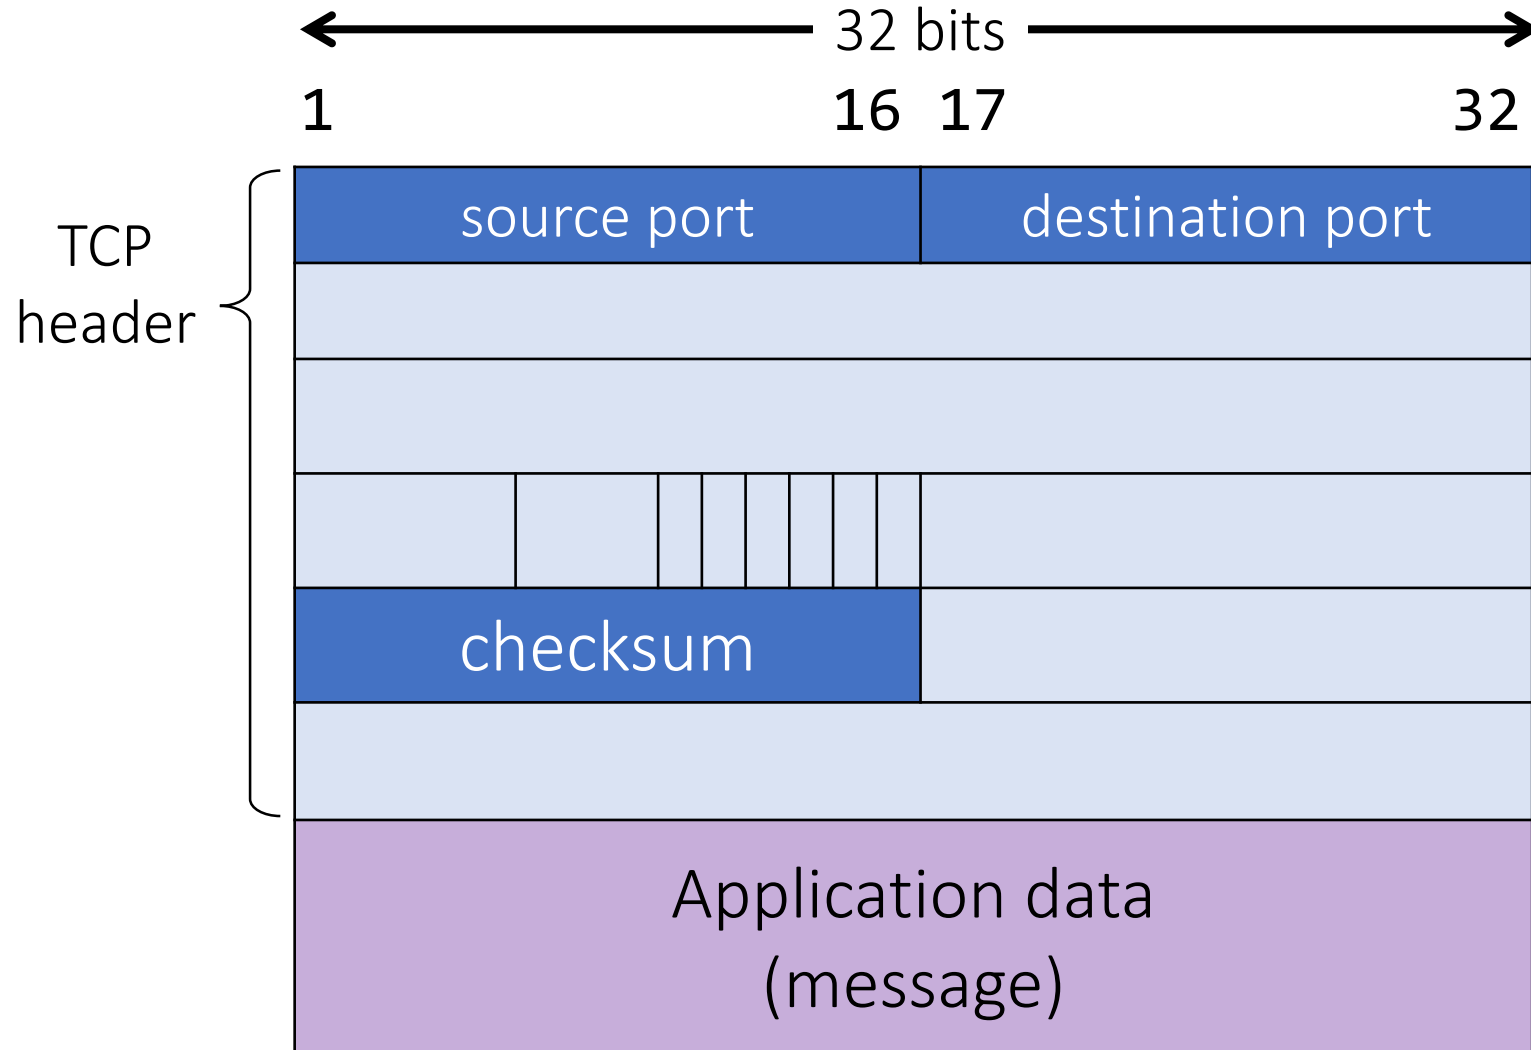
- (srcIPAddr, srcPort, destIPAddr, destPort)
- Receiver uses all four values to direct a segment to the appropriate socket.



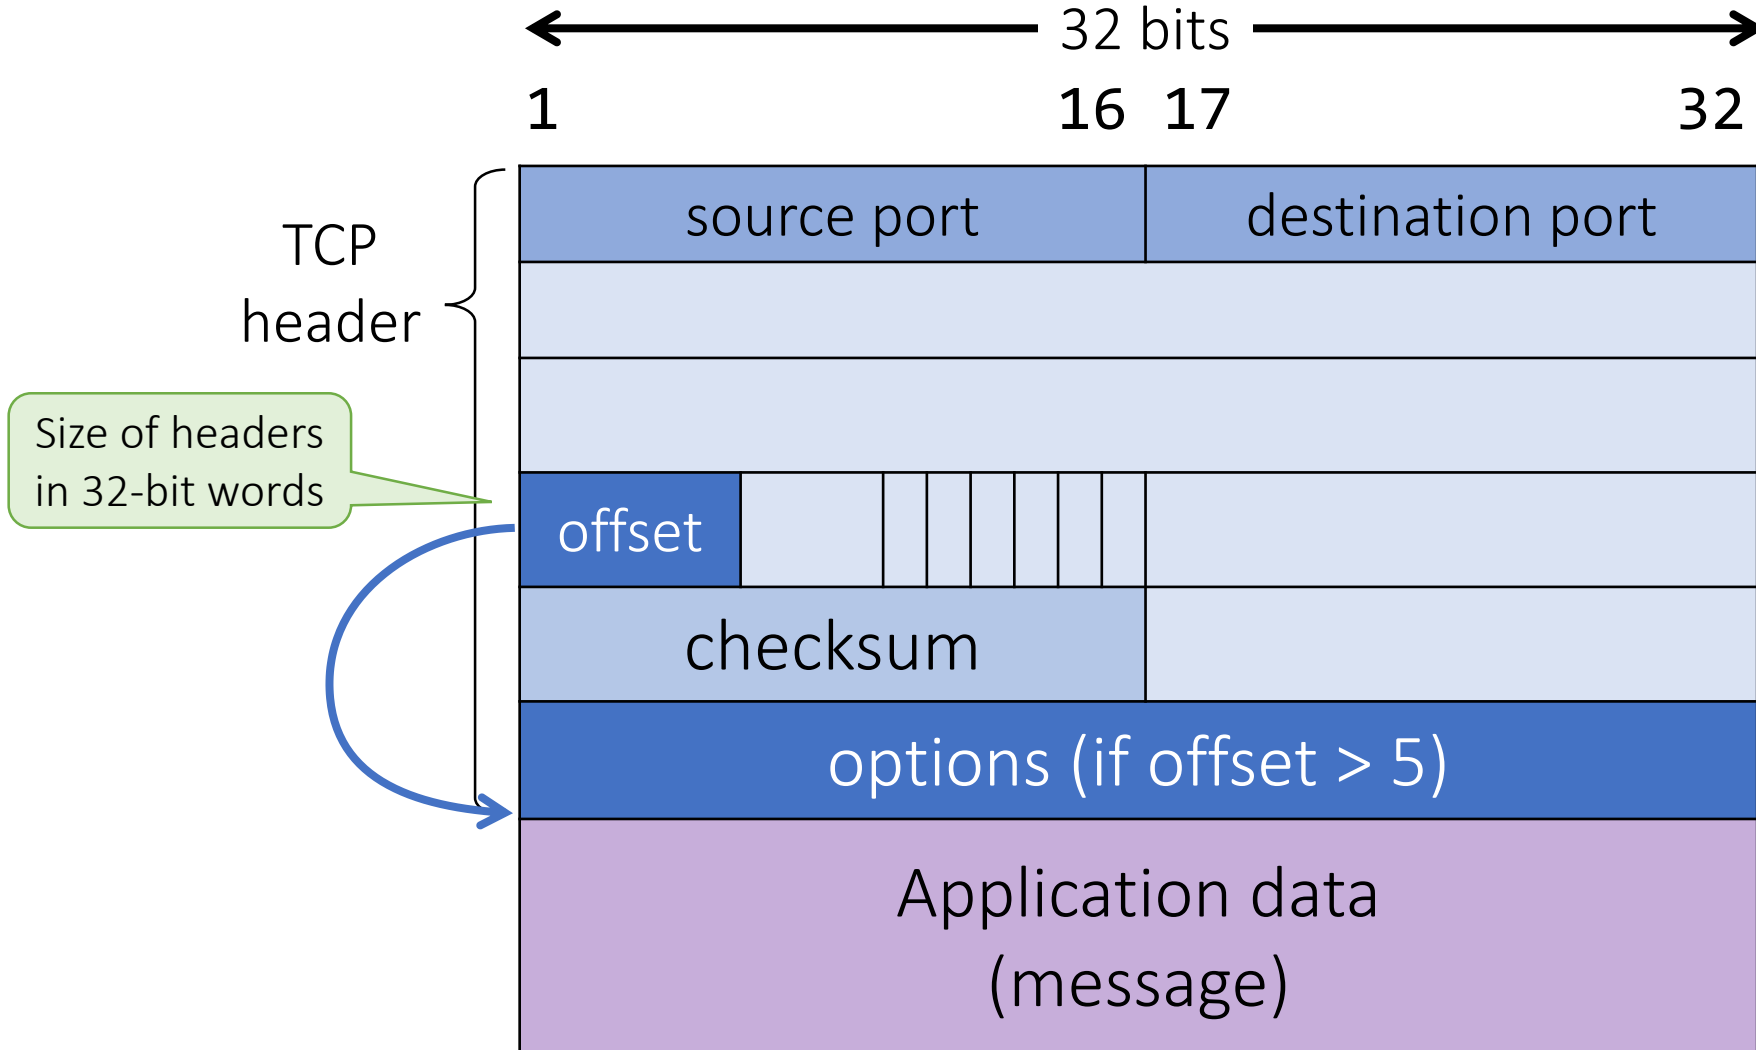
# TCP segment structure



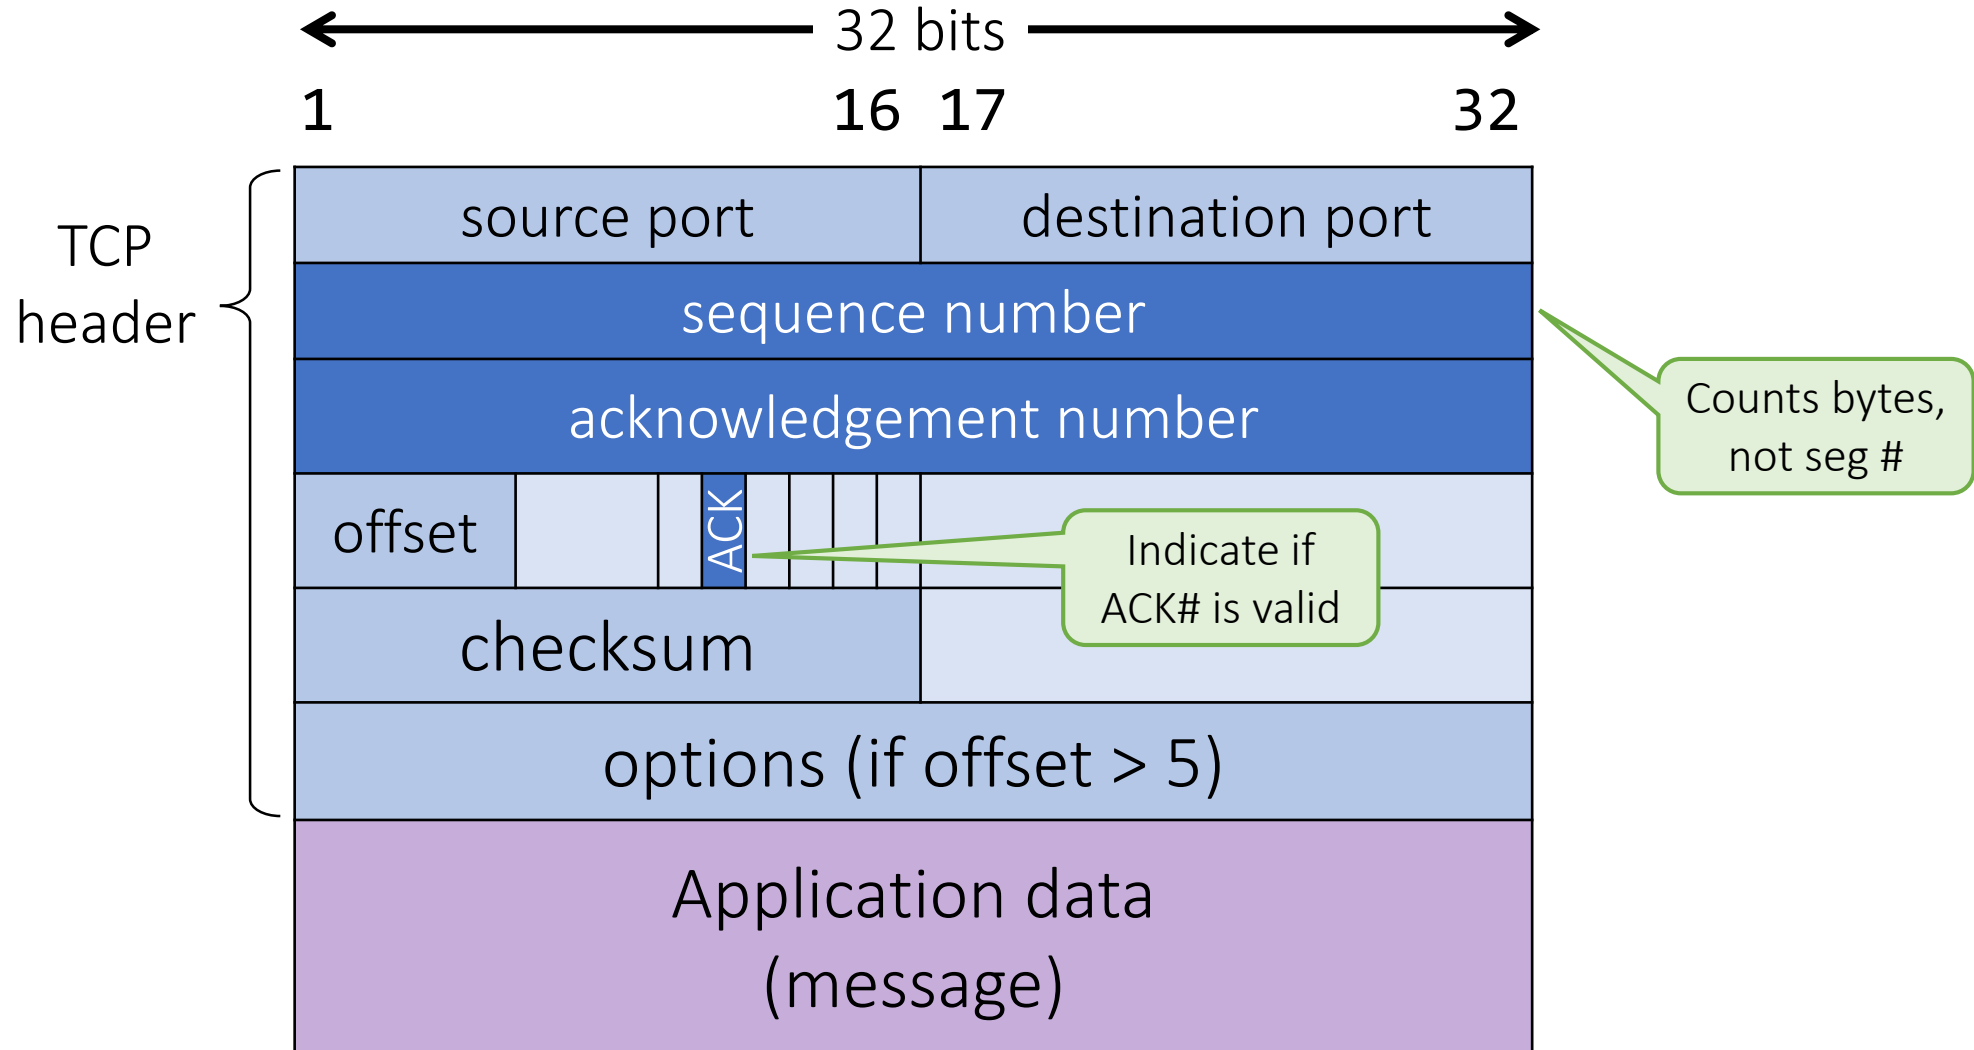
# TCP segment structure



# TCP segment structure



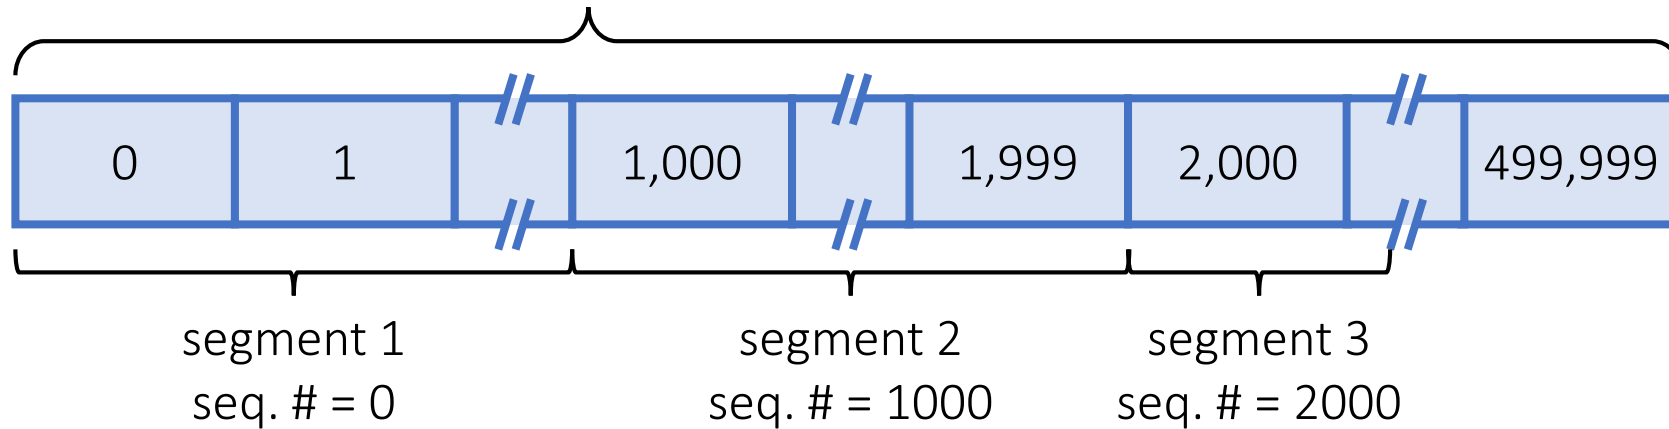
# TCP Sequence numbers



# TCP Sequence Number

“byte number” of the first byte of data in the segment

Example: 500,000-byte file with 1,000-byte segments






# TCP Acknowledgement Number

Seq. # of the next byte of data expected

TCP only ACKs up to the missing byte (cumulative ACK)

Receives 1,000 byte seq. # 5,000



Receive window/buffer


Sends ACK # 6,000

# TCP Acknowledgement Number

Seq. # of the next byte of data expected

TCP only ACKs up to the missing byte (cumulative ACK)

Receives 500 byte seq. # 6,500



Sends ACK # 6,000


Note: TCP specs does not say how out-of-order segments should be handled

# TCP Acknowledgement Number

Seq. # of the next byte of data expected

TCP only ACKs up to the missing byte (cumulative ACK)

Receives 500 byte seq. # 6,000



Receive window/buffer

Sends ACK # 7,000

# TCP Acknowledgement Number

Seq. # of the next byte of data expected

TCP only ACKs up to the missing byte (cumulative ACK)

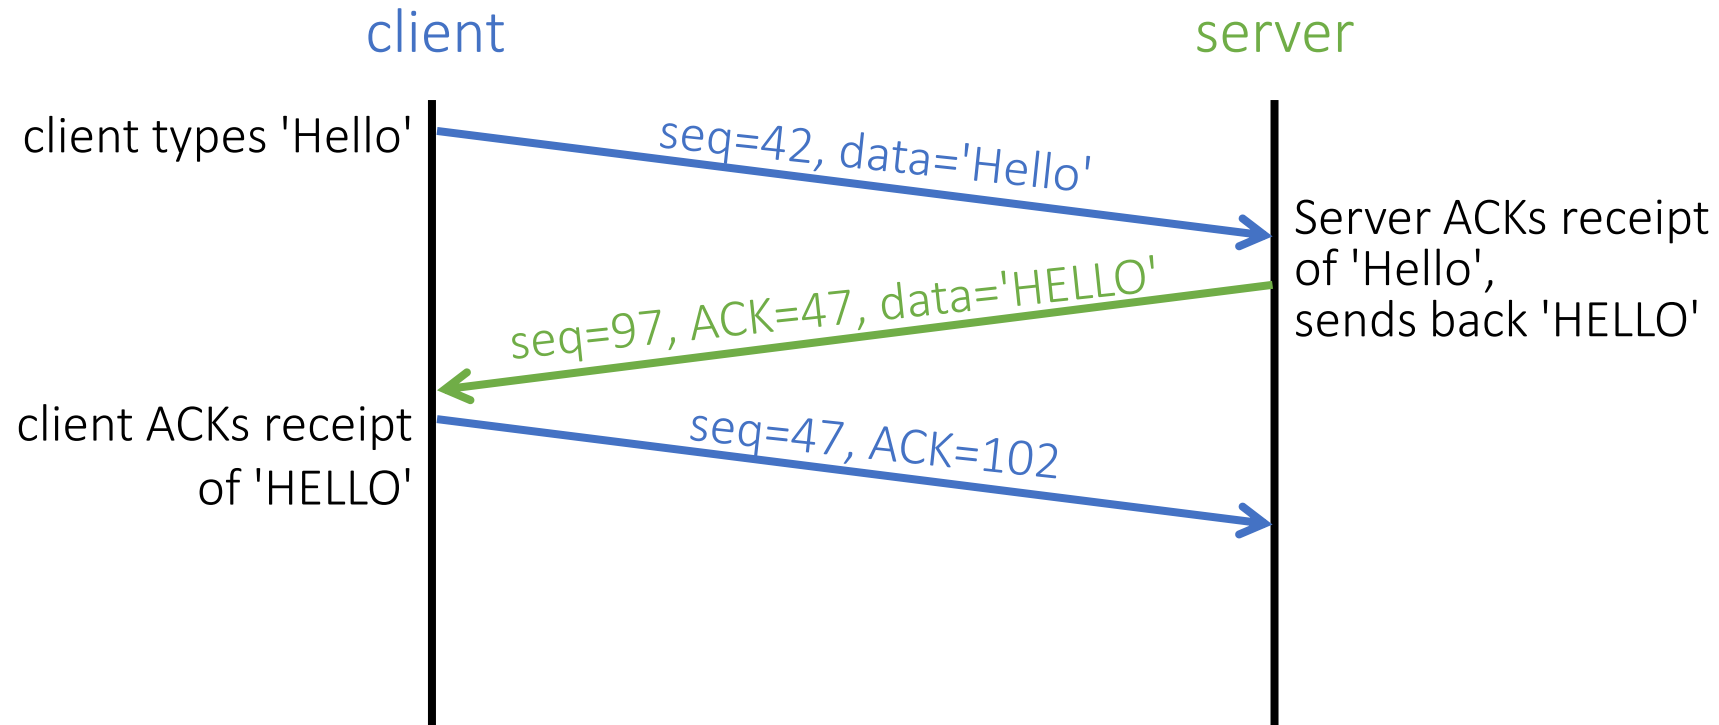
Receives 500 byte seq. # 6,000



Sends ACK # 7,000

# TCP Echo Server

TCP is full duplex, i.e., bi-directional data flow



ACKs are “piggybacked” on data segment

Question: How large should a segment be?

Maximum Segment Size (MSS)

typically derived from link-layer's MTU

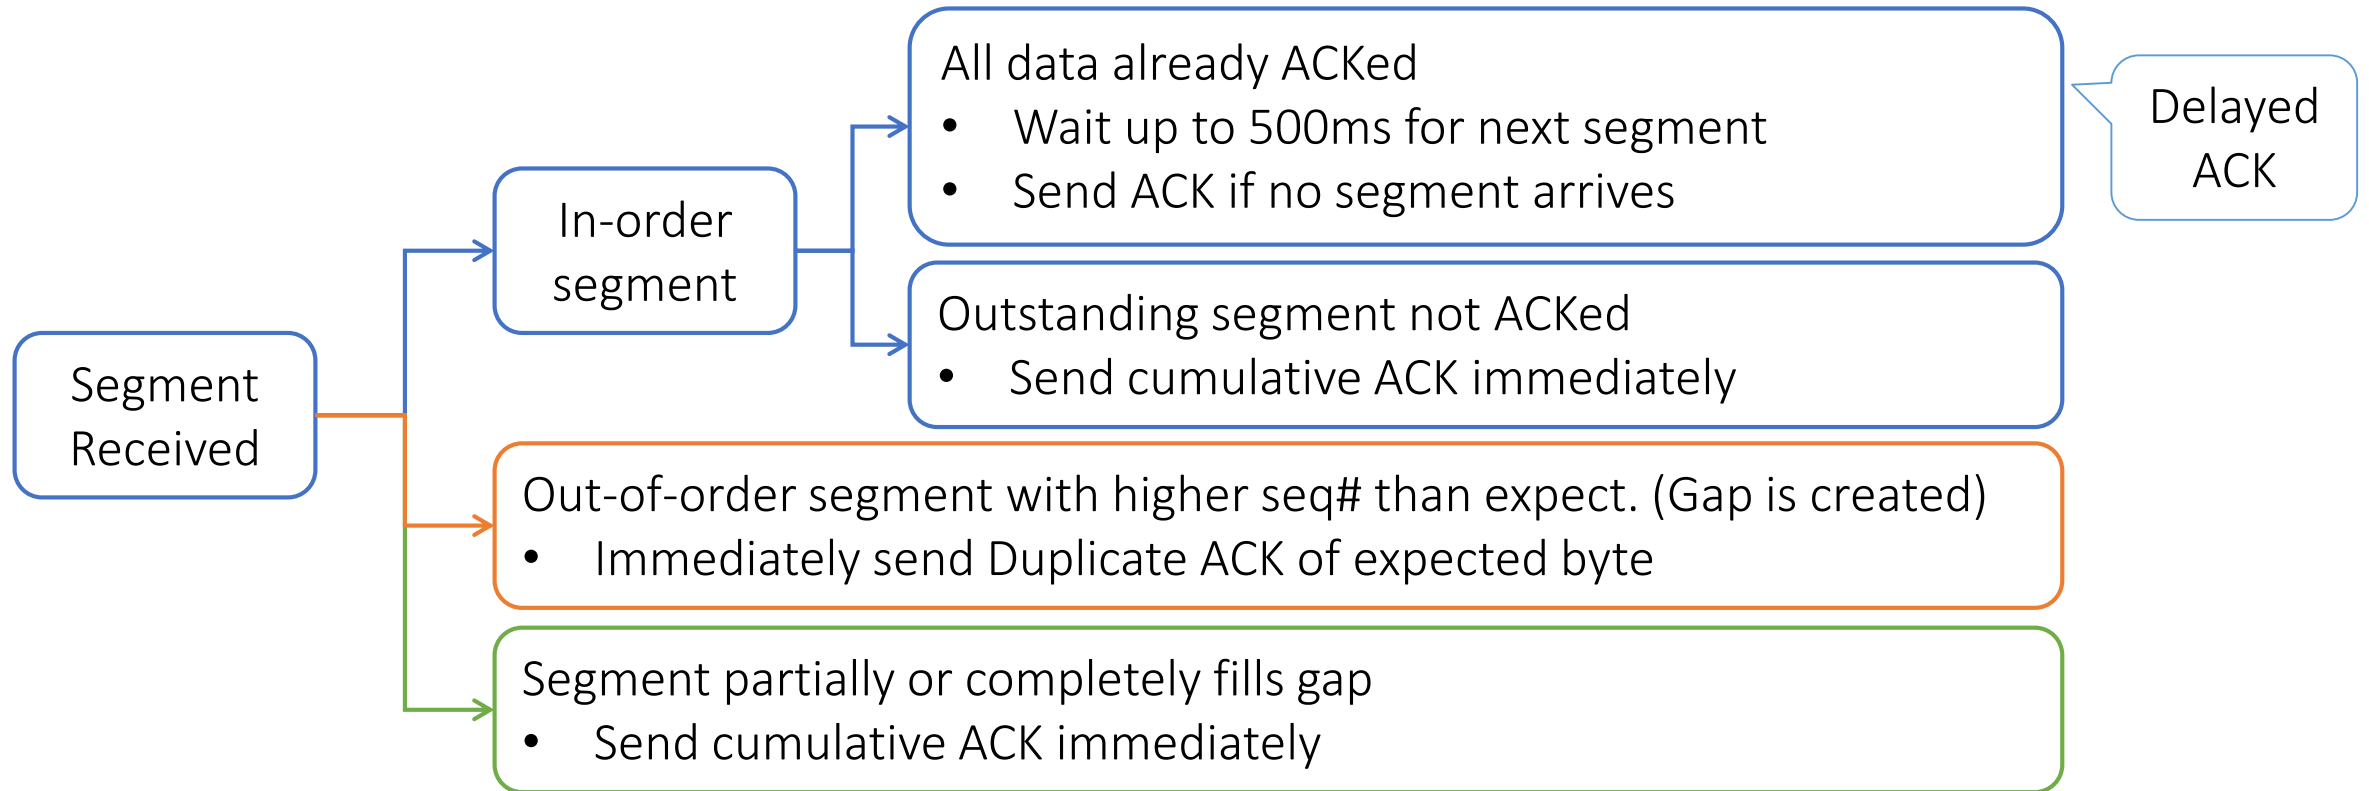
# TCP Sender Events (simplified)

```
loop(forever)
  switch(event)
    event: data received from application
      create TCP segment with nextSeqNum
      if (timer not currently running) ← Sender only keeps one timer
        start timer
      pass segment to IP
      nextSeqNum += length(data)

    event: timer timeout
      retransmit unacknowledged segment with smallest seq num ← Retransmit only oldest unACK segment
      start timer



    event: ACK received, with ACK num #y ← Cumulative ACK
      if (y > sendBase)
        sendBase = y
        if (there are still unacknowledged segments)
          start timer
```

# TCP Receiver Events

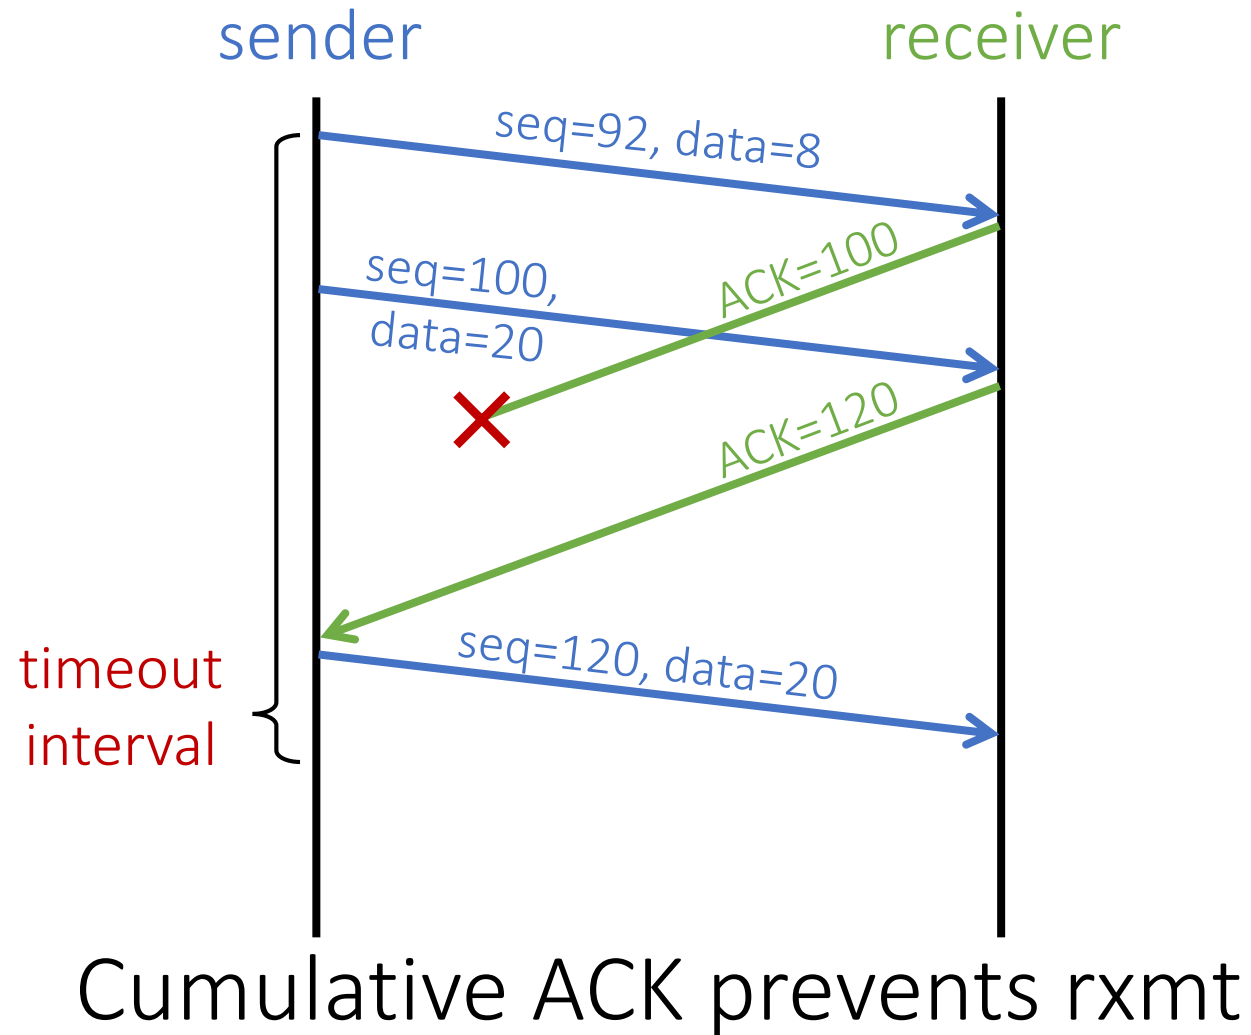




# TCP Timeout/Retransmission



# TCP Timeout/Retransmission



# TCP Timeout Value

Too short?


- Premature timeout and unnecessary retransmissions.

Too long?

- Slow reaction to loss

Key point:  $\text{Timeout} > \text{RTT}$

- But RTT can vary!



# Estimating RTT

Take SampleRTT ( $RTT_s$ )

- *Once every RTT*

Compute EstimatedRTT ( $RTT_\epsilon$ )

- $RTT_\epsilon = (1 - \alpha) \cdot RTT_\epsilon + \alpha \cdot RTT_s$
- typical value of  $\alpha = \frac{1}{8}$

Exponential Weighted Moving Average (EWMA)

# Setting Retransmission Time Out (RTO)

Compute Deviation of RTT

$$- RTT_{dev} = (1 - \beta) \cdot RTT_{dev} + \beta \cdot |RTT_s - RTT_\varepsilon|$$

- typical value of  $\beta = \frac{1}{4}$

RTO Interval is set to

$$RTT_\varepsilon + 4 \times RTT_{dev}$$


estimated RTT

“safety margin”

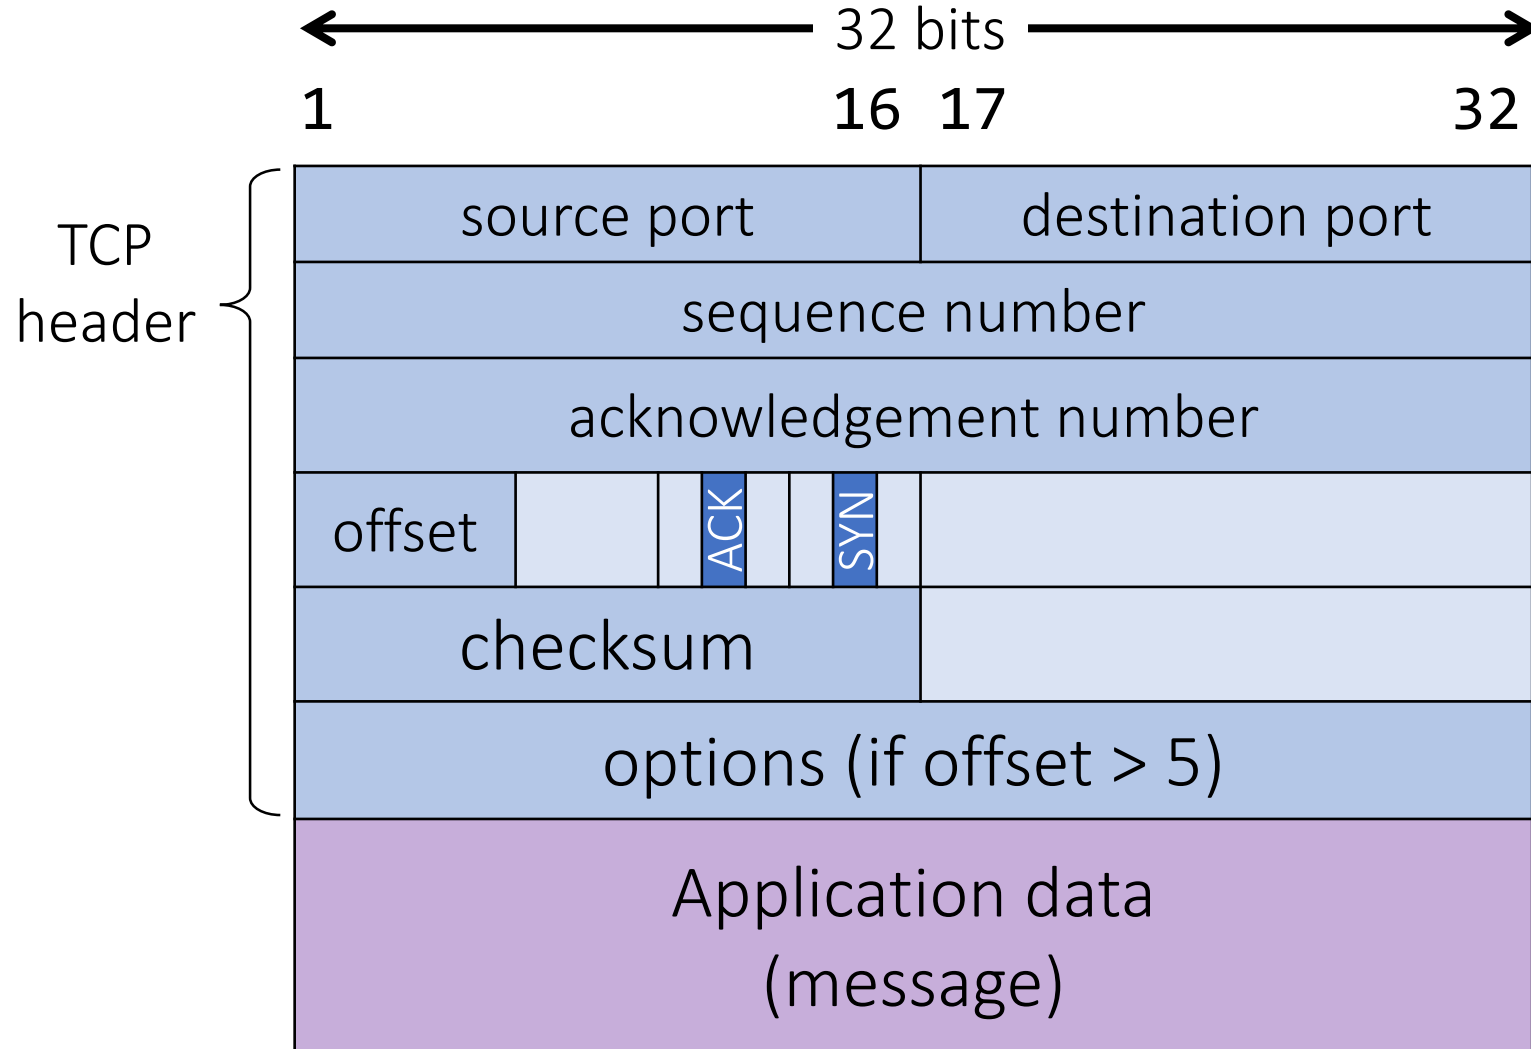
# TCP Fast Retransmission [RFC 2001]

Timeout is often relatively long

- If 3 Duplicate ACKs are received, resend segment immediately



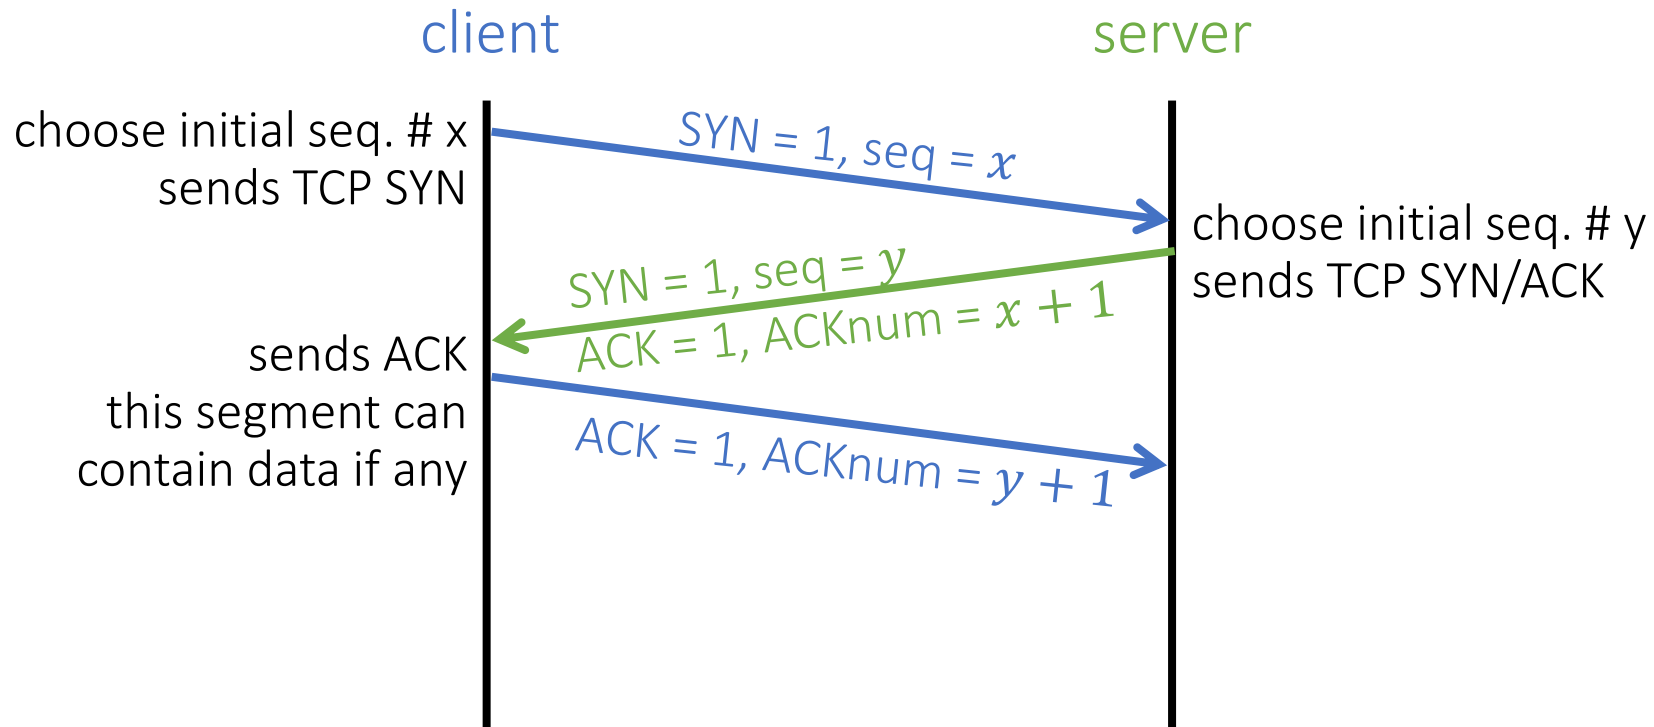
# TCP Connection Establishment



# TCP Connection Establishment

Established using a 3-way handshake

- Agree on connection and exchange parameters





# Half-open Connections

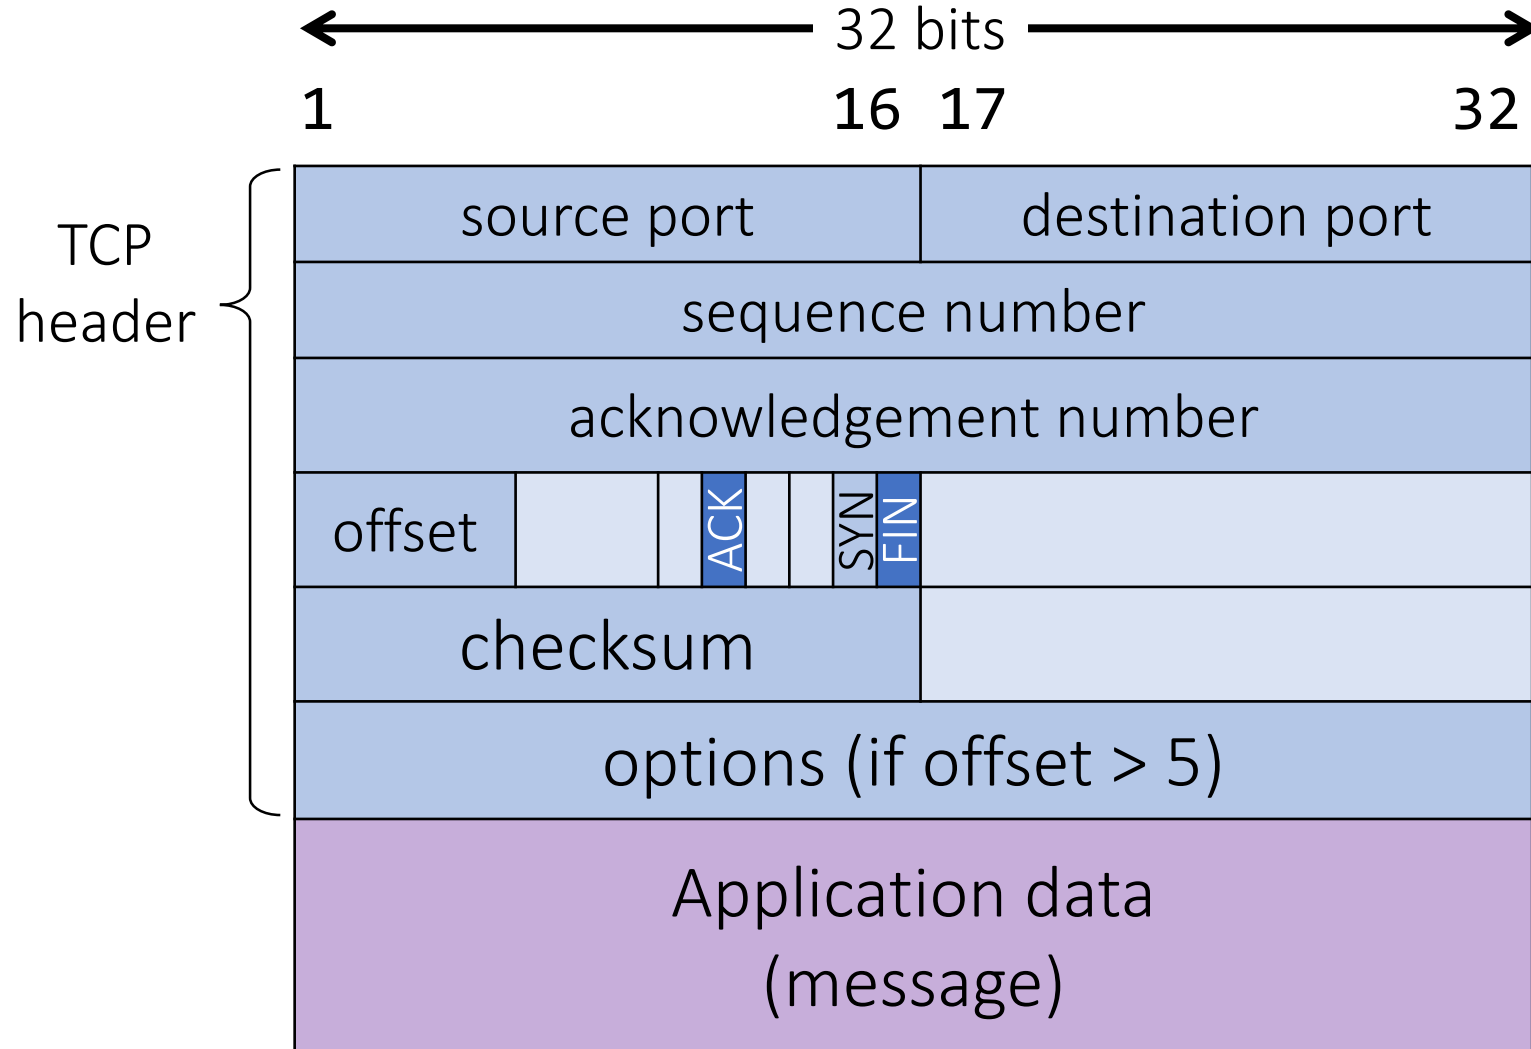
## SYN Flooding

- DoS style attack by sending SYN

## SYN/ACK Flooding

- DoS style to overwhelm network

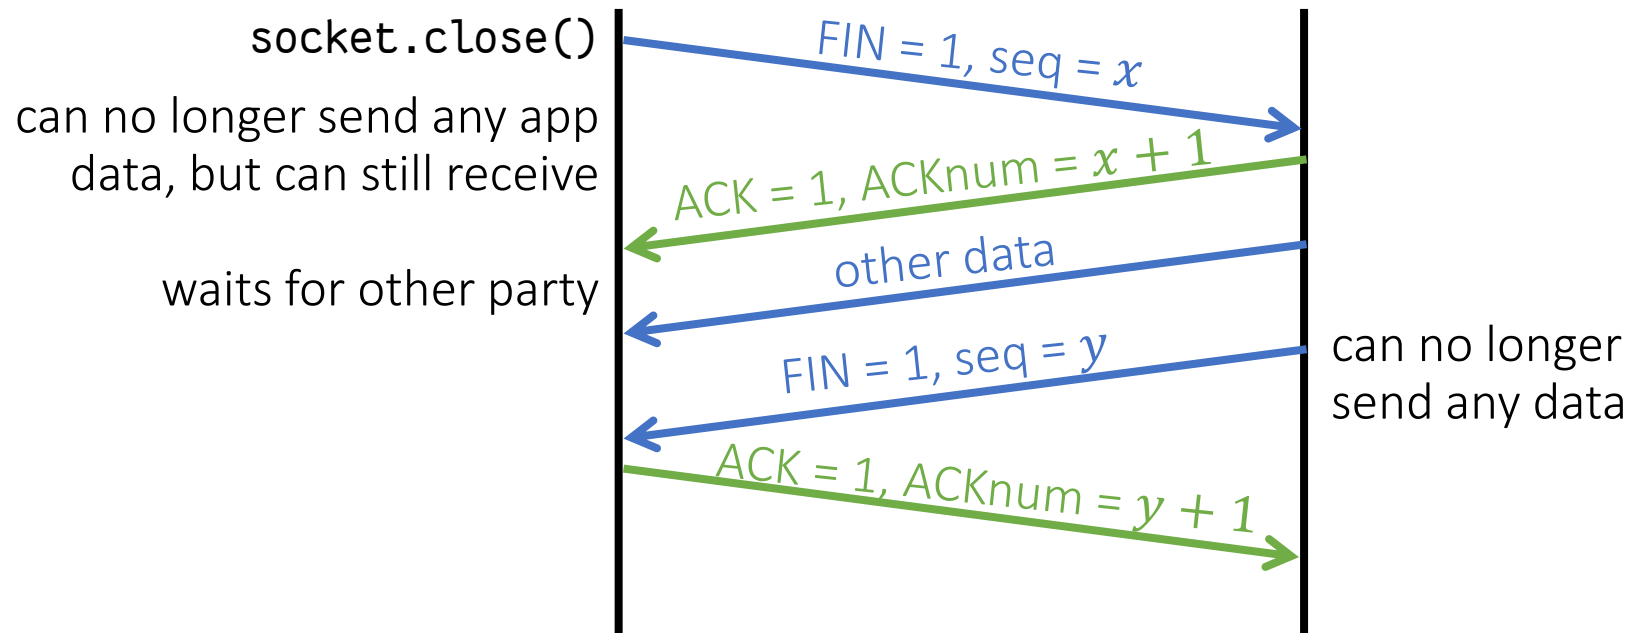
# TCP Closing Connection



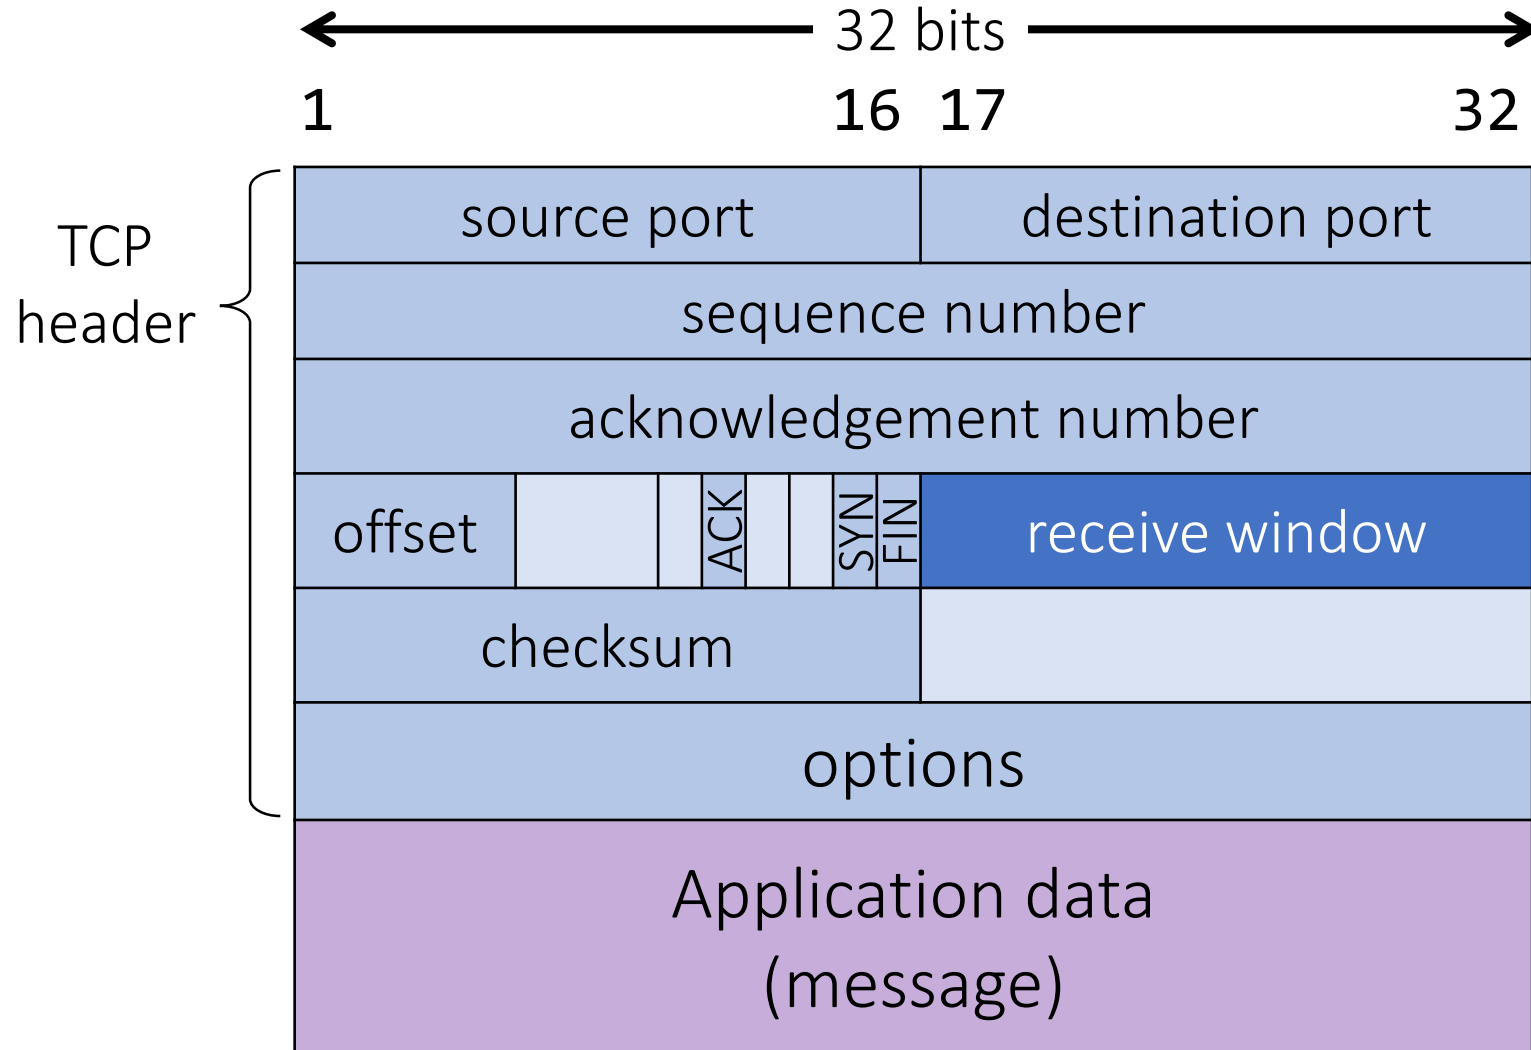
# TCP Closing Connection

Each side closes their own side of the connection

- Send segments with FIN bit set
- No more sending of data after FIN, but can still receive

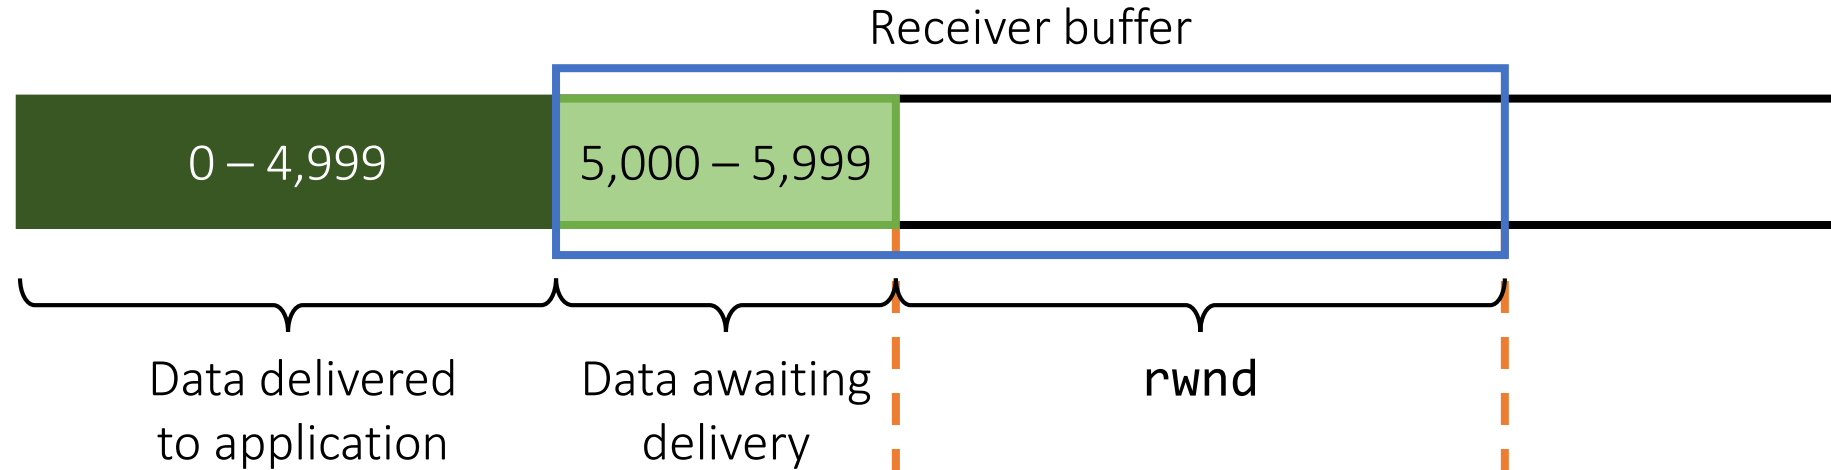


# TCP Flow Control

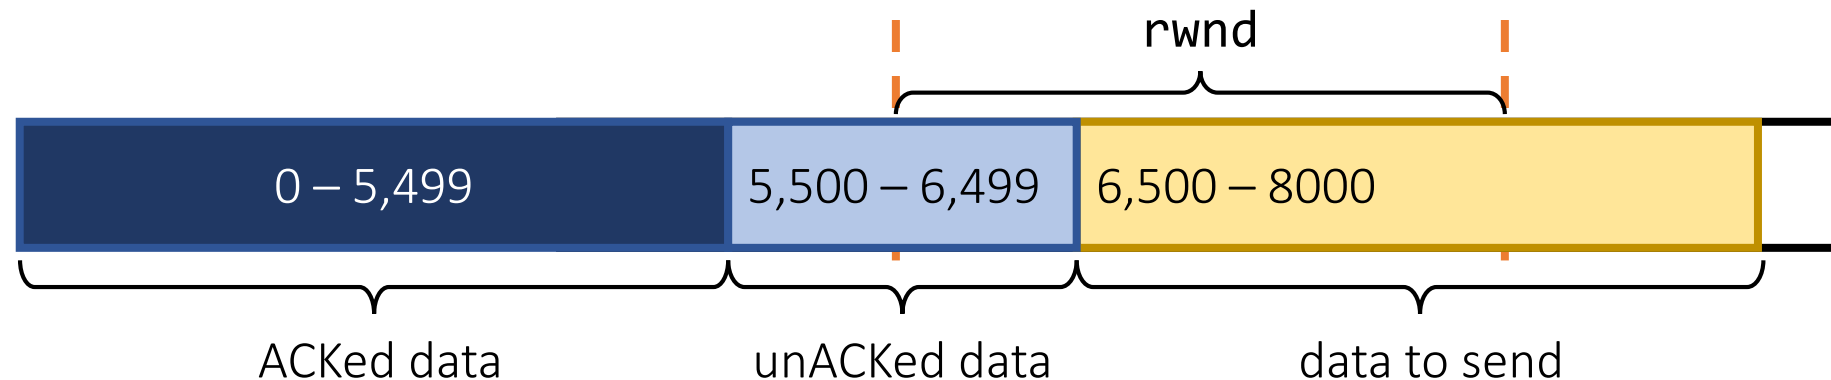


# TCP Flow Control

Receiver buffers data to application

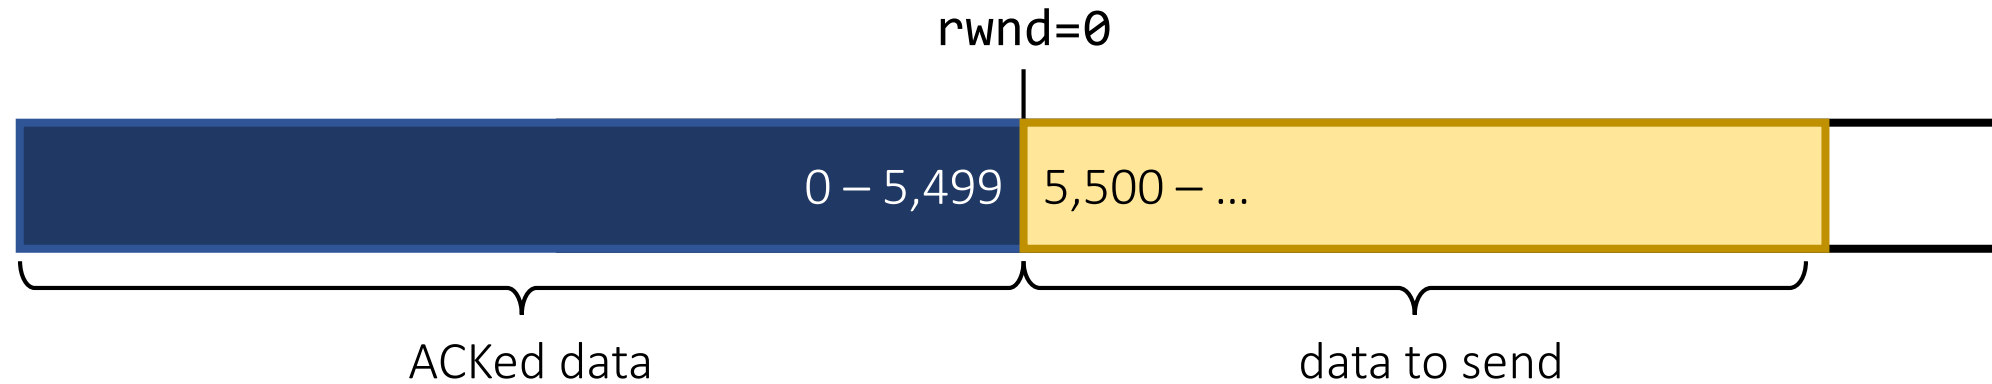


Tells sender how much data it can send



# TCP Flow Control

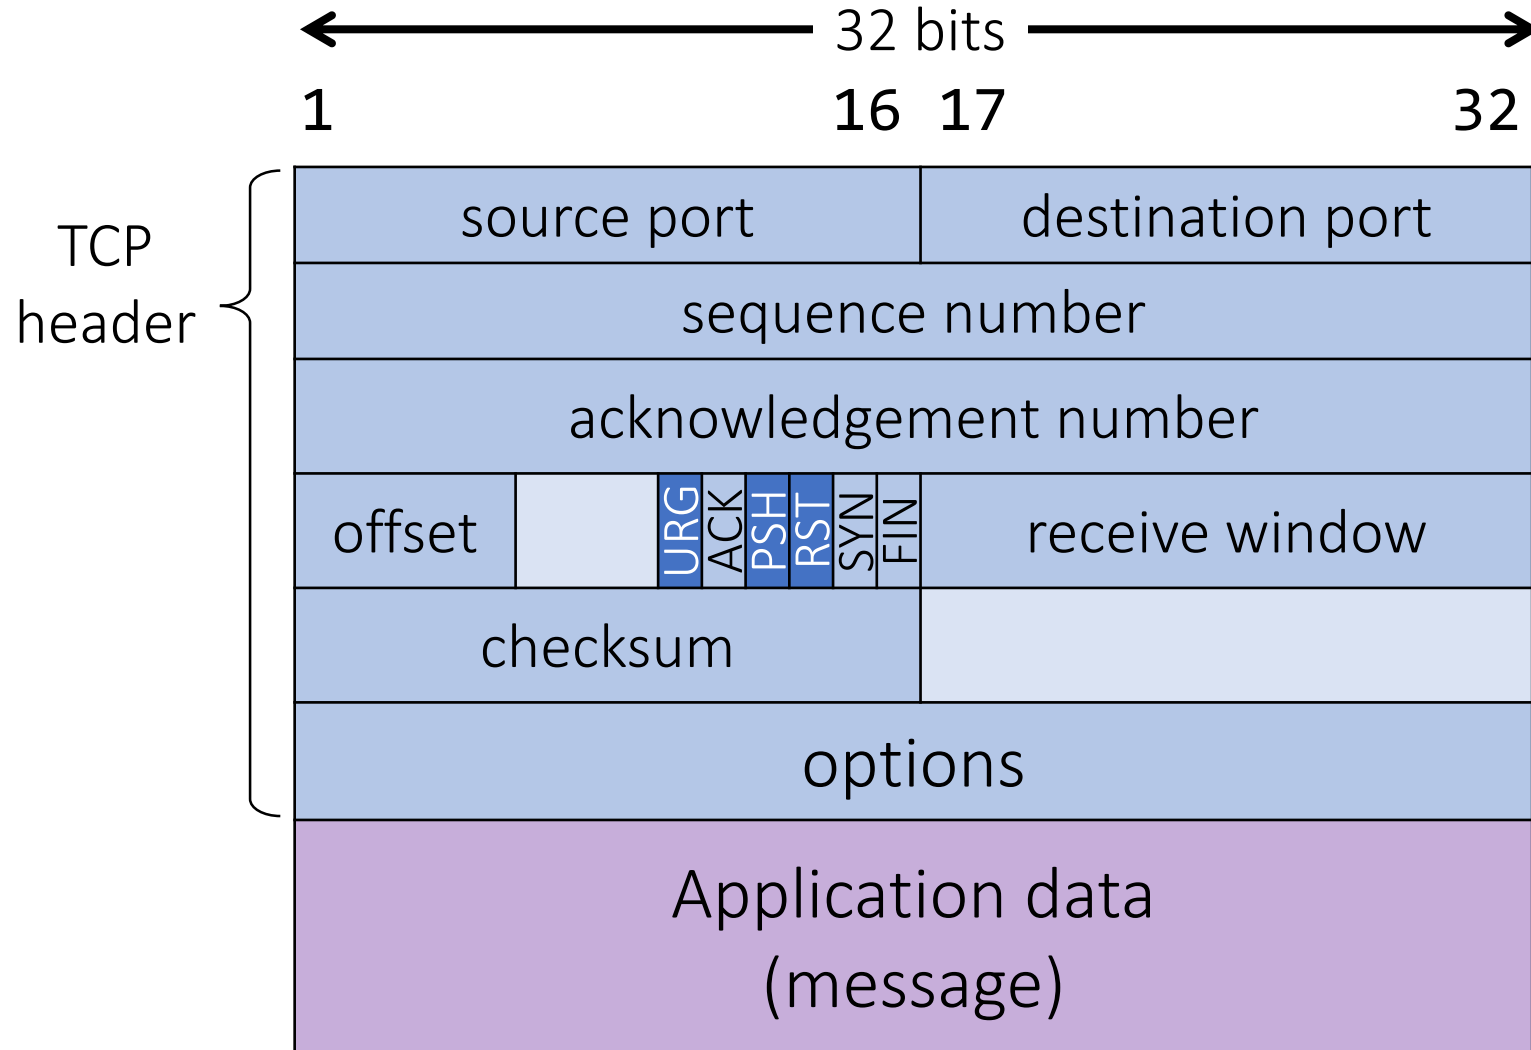
What if rwnd is 0, i.e., full?



How does the sender know when it empties?

- Sender sends a 0-data segment
- a.k.a zero-window probe

# TCP segment structure



# What we did not cover

## TCP congestion control (Chapter 3.6 & 3.7)

- Be polite and send less if network is congested

## Other TCP options

- e.g., SACK, Timestamp

## Other TCP techniques

- e.g., FACK, DSACK



# Summary

## Connectionless: UDP

- Segment structure
- Computing Checksum

## Connection-oriented: TCP

- Segment structure
- Reliable data transfer
- Setting and updating RTO
- Establishing and closing connection
- Flow control